



Postgres Plus® Advanced Server Performance and Scalability Guide

Postgres Plus Advanced Server 9.2

November 8, 2012

Postgres Plus Advanced Server Performance Features Guide, Version 3.0
by EnterpriseDB Corporation
Copyright © 2012 EnterpriseDB Corporation. All rights reserved.

Table of Contents

1	Introduction	6
1.1	Typographical Conventions Used in this Guide	7
2	Infinite Cache	8
2.1	Installing Infinite Cache	12
2.2	Configuring the Infinite Cache Server	14
2.2.1	Modifying Infinite Cache Settings	14
2.2.2	Enabling Infinite Cache	14
2.2.3	Controlling the Infinite Cache Server	16
2.3	Dynamically Modifying Infinite Cache Server Nodes	18
2.4	Controlling the edb-icache Daemons	19
2.4.1	Command Line Options	20
2.4.2	edb-icache-tool	21
2.5	Warming the edb-icache Servers	23
2.5.1	The edb_icache_warm() Function	23
2.5.2	Using the edb_icache_warm Utility	24
2.6	Retrieving Statistics from Infinite Cache	25
2.6.1	Using edb_icache_stats()	25
2.6.2	edb_icache_server_list	26
2.7	Retrieving Table Statistics	27
2.7.1	pg_statio_all_tables	28
2.7.2	pg_statio_sys_tables	29
2.7.3	pg_statio_user_tables	30
2.7.4	pg_statio_all_indexes	30
2.7.5	pg_statio_sys_indexes	31
2.7.6	pg_statio_user_indexes	31
2.8	edb_icache_server_enable()	32
2.9	Infinite Cache Log Entries	32
2.10	Allocating Memory to the Cache Servers	33
3	Dynatune	34
3.1.1	edb_dynatune	34
3.1.2	edb_dynatune_profile	35
4	Dynamic Runtime Instrumentation Tools Architecture (DRITA)	36

4.1.1	Initialization Parameters.....	36
4.1.2	Setting up and Using DRITA.....	36
4.2	DRITA Functions	36
4.2.1	get_snaps().....	36
4.2.2	sys_rpt().....	37
4.2.3	sess_rpt().....	38
4.2.4	sessid_rpt().....	39
4.2.5	sesshist_rpt().....	40
4.2.6	purgesnap()	41
4.2.7	truncsnap().....	42
4.3	Simulating Statspack AWR Reports.....	43
4.3.1	edbreport().....	43
4.3.2	stat_db_rpt().....	56
4.3.3	stat_tables_rpt().....	56
4.3.4	statio_tables_rpt().....	58
4.3.5	stat_indexes_rpt().....	59
4.3.6	statio_indexes_rpt().....	61
4.4	Performance Tuning Recommendations	63
4.5	Event Descriptions.....	64
4.6	Catalog Views.....	66
4.6.1	edb\$system_waits.....	66
4.6.2	edb\$session_waits	66
4.6.3	edb\$session_wait_history	67
5	DBMS_PROFILER	68
5.1	Querying the DBMS_PROFILER Tables and View	70
5.2	DBMS_PROFILER Functions and Procedures	78
5.2.1	FLUSH_DATA	78
5.2.2	GET_VERSION	79
5.2.3	INTERNAL_VERSION_CHECK	79
5.2.4	PAUSE_PROFILER	79
5.2.5	RESUME_PROFILER.....	80
5.2.6	START_PROFILER	80
5.2.7	STOP_PROFILER	81
5.3	DBMS_PROFILER - Reference.....	82

5.3.1	PLSQL_PROFILER_RUNS.....	82
5.3.2	PLSQL_PROFILER_UNITS.....	82
5.3.3	PLSQL_PROFILER_DATA	83
5.3.4	PLSQL_PROFILER_RAWDATA	83
6	Index Advisor.....	88
6.1	Index Advisor Components.....	89
6.2	Index Advisor Configuration	90
6.3	Using Index Advisor.....	92
6.3.1	Using the pg_advise_index Utility	93
6.3.2	Using Index Advisor at the psql Command Line	94
6.4	Reviewing the Index Advisor Recommendations	95
6.4.1	Using the show_index_recommendations() Function.....	96
6.4.2	Querying the index_advisor_log Table.....	96
6.4.3	Querying the index_recommendations View	98
6.5	Limitations	100
7	Other Performance Features	101
7.1	SQL Profiler	101
7.2	Query Optimization Hints	101
7.3	Hi-Speed Bulk Loader.....	102
7.4	Bulk Collect Fetch and Binding.....	103
7.5	Multi-Threaded Replication.....	103

1 Introduction

This guide describes the performance features found in Postgres Plus Advanced Server.

- Infinite Cache allows you to utilize memory on other computers connected to your network to increase the amount of memory in the shared buffer cache.
- Dynatune makes optimal use of the system resources that are available on the host machine.
- The Dynamic Runtime Instrumentation Tools Architecture (DRITA) records *wait events* that affect system performance; DRITA also offers a set of tools for inspecting those events.
- Optimizer Hints are directives that you embed in comment-like syntax immediately following the `SELECT`, `UPDATE`, or `DELETE` key words to influence the query optimizer.

1.1 *Typographical Conventions Used in this Guide*

Certain typographical conventions are used in this manual to clarify the meaning and usage of various commands, statements, programs, examples, etc. This section provides a summary of these conventions.

In the following descriptions a *term* refers to any word or group of words that are language keywords, user-supplied values, literals, etc. A term's exact meaning depends upon the context in which it is used.

- *Italic font* introduces a new term, typically, in the sentence that defines it for the first time.
- Fixed-width (mono-spaced) font is used for terms that must be given literally such as SQL commands, specific table and column names used in the examples, programming language keywords, etc. For example, `SELECT * FROM emp;`
- *Italic fixed-width font* is used for terms for which the user must substitute values in actual usage. For example, `DELETE FROM table_name;`
- A vertical pipe | denotes a choice between the terms on either side of the pipe. A vertical pipe is used to separate two or more alternative terms within square brackets (optional choices) or braces (one mandatory choice).
- Square brackets [] denote that one or none of the enclosed term(s) may be substituted. For example, [a | b], means choose one of “a” or “b” or neither of the two.
- Braces { } denote that exactly one of the enclosed alternatives must be specified. For example, { a | b }, means exactly one of “a” or “b” must be specified.
- Ellipses ... denote that the proceeding term may be repeated. For example, [a | b] ... means that you may have the sequence, “b a a b a”.

2 Infinite Cache

Database performance is typically governed by two competing factors:

- Memory access is fast; disk access is slow.
- Memory space is scarce; disk space is abundant.

Postgres Plus Advanced Server tries very hard to minimize disk I/O by keeping frequently used data in memory. When the first server process starts, it creates an in-memory data structure known as the *buffer cache*. The buffer cache is organized as a collection of 8K (8192 byte) pages: each page in the buffer cache corresponds to a page in some table or index. The buffer cache is shared between all processes servicing a given database.

When you select a row from a table, Advanced Server reads the page that contains the row into the shared buffer cache. If there isn't enough free space in the cache, Advanced Server *evicts* some other page from the cache. If Advanced Server evicts a page that has been modified, that data is written back out to disk; otherwise, it is simply discarded. Index pages are cached in the shared buffer cache as well.

Figure 1.1 demonstrates the flow of data in a typical Advanced Server session:

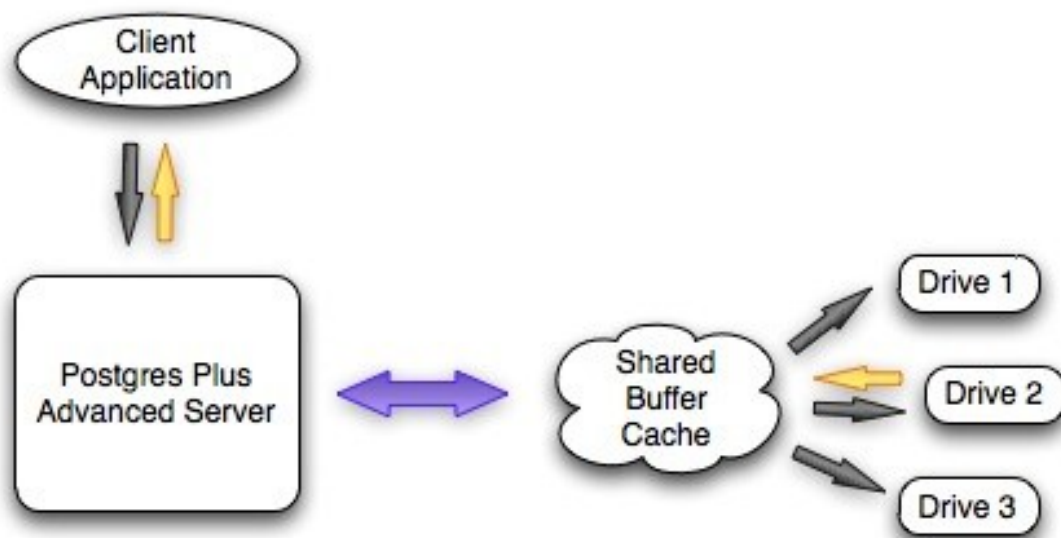


Figure 1.1 – Data Flow

A client application sends a query to the Postgres server and the server searches the shared buffer cache for the required data. If the requested data is found in the cache, the server immediately sends the data back to the client. If not, the server reads the page that holds the data into the shared buffer cache, evicting one or more pages if necessary. If the server decides to evict a page that has been modified, that page is written to disk.

As you can see, a query will execute much faster if the required data is found in the shared buffer cache.

One way to improve performance is to increase the amount of memory that you can devote to the shared buffer cache. However, most computers impose a strict limit on the amount of RAM that you can install. To help circumvent this limit, Infinite Cache lets you utilize memory from other computers connected to your network.

With Infinite Cache properly configured, Advanced Server will dedicate a portion of the memory installed on each *cache server* as a secondary memory cache. When a client application sends a query to the server, the server first searches the shared buffer cache for the required data; if the requested data is not found in the cache, the server searches for the necessary page in one of the cache servers.

Figure 1.2 shows the flow of data in an Advanced Server session with Infinite Cache:

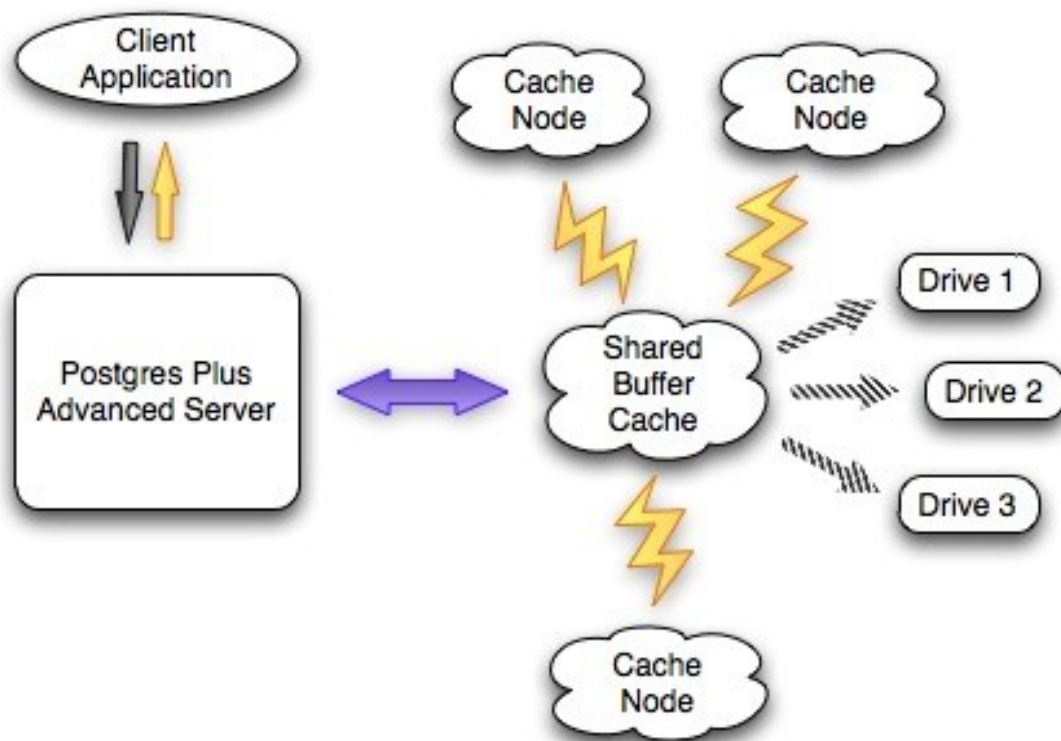


Figure 1.2 – Data flow with Infinite Cache

When a client application sends a query to the server, the server searches the shared buffer cache for the required data. If the requested data is found in the cache, the server immediately sends the data back to the client. If not, the server sends a request for the page to a specific cache server; if the cache server holds a copy of the page it sends the data back to the server and the server copies the page into the shared buffer cache. If the required page is not found in the primary cache (the shared buffer cache) or in the secondary cache (the cloud of cache servers), Advanced Server must read the page from disk. Infinite Cache improves performance by utilizing RAM from other computers on your network in order to avoid reading frequently accessed data from disk.

Updating the Cache Node Configuration

You can add or remove cache servers without restarting the database server by adding or deleting cache nodes from the list defined in the `edb_icache_servers` configuration parameter. For more information about changing the configuration parameter, see Section 2.2.2.

When you add one or more cache nodes, the server re-allocates the cache, dividing the cache evenly amongst the servers; each of the existing cache servers loses a percentage of the information that they have cached. You can calculate the percentage of the cache that remains valid with the following formula:

$$(existing_nodes * 100) / (existing_nodes + new_nodes)$$

For example, if an Advanced Server installation with three existing cache nodes adds an additional cache node, 75% of the existing cache remains valid after the reconfiguration.

If cache nodes are removed from a server, the data that has been stored on the remaining cache nodes is preserved. If one cache server is removed from a set of five cache servers, Advanced Server preserves the 80% of the distributed cache that is stored on the four remaining cache nodes.

When you change the cache server configuration (by adding or removing cache servers), the portion of the cache configuration that is preserved is not re-written unless the cache is completely re-warmed using the `edb_icache_warm()` function or `edb_icache_warm` utility. If you do not re-warm the cache servers, new cache servers will accrue cache data as queries are performed on the server.

Infinite Cache Offers a Second Performance Advantage: Compression.

Without Infinite Cache, Advanced Server will read each page from disk as an 8K chunk; when a page resides in the shared buffer cache, it consumes 8K of RAM. With Infinite Cache, Postgres can *compress* each page before sending it to a cache server. A compressed page can take significantly less room in the secondary cache, making more space available for other data and effectively increasing the size of the cache. A

compressed page consumes less network bandwidth as well, decreasing the amount of time required to retrieve a page from the secondary cache.

The fact that Infinite Cache can compress each page may make it attractive to configure a secondary cache server on the same computer that runs your Postgres server. If, for example, your computer is configured with 6GB of RAM, you may want to allocate a smaller amount (say 1GB) for the primary cache (the shared buffer cache) and a larger amount (4GB) to the secondary cache (Infinite Cache), reserving 1GB for the operating system. Since the secondary cache resides on the same computer, there is very little overhead involved in moving data between the primary and secondary cache. All data stored in the Infinite Cache is compressed so the secondary cache can hold many more pages than would fit into the (uncompressed) shared buffer cache. If you had allocated 5GB to the shared buffer cache, the cache could hold no more than 65000 pages (approximately). By assigning 4GB of memory to Infinite Cache, the cache may be able to hold 130000 pages (at 2x compression), 195000 pages (at 3x compression) or more. The compression factor that you achieve is determined by the amount of redundancy in the data itself and the `edb_icache_compression_level` parameter.

To use Infinite Cache, you must specify a list of one or more cache servers (computers on your network) and start the `edb_icache` daemon on each of those servers.

Infinite Cache is supported on Linux, HP-UX and Solaris systems only.

Please Note: Infinite Cache and the `effective_io_concurrency` parameter can potentially interfere with each other. You should disable asynchronous I/O requests (by setting the value of `effective_io_concurrency` to 0 in the `postgresql.conf` file) if you enable the Infinite Cache feature.

2.1 Installing Infinite Cache

Postgres Plus Advanced Server includes Infinite Cache functionality as part of a standard installation on a Linux, HP-UX or Solaris system. The Advanced Server installation wizard can optionally install only the Infinite Cache daemon on supporting cache servers without installing Advanced Server.

To install Advanced Server with Infinite Cache functionality, confirm that the box next to the `Database Server` option (located on the `Setup: Select Components` window, shown in Figure 1.3) is selected when running the installation wizard.

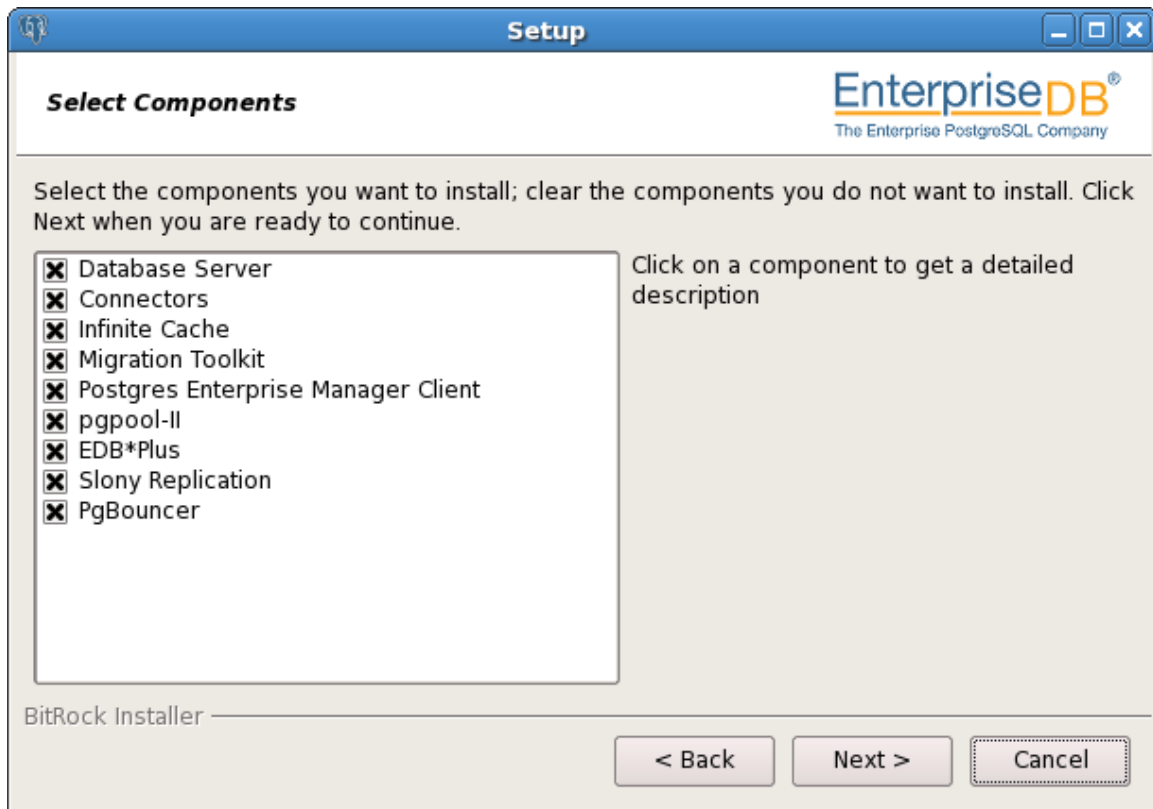


Figure 1.3: The Setup: Select Components window.

Selecting the `Database Server` option installs the following Infinite Cache components:

- The `ppas-infinitecache-9.2` service script.
- The Infinite Cache configuration file (`ppas-infinitecache-9.2`).
- A command line tool that allows you to pre-load the cache servers (`edb-icache-warm`).
- The `edb_icache` libraries (code libraries required by the `edb-icache` daemon).

The installation wizard can selectively install only the Infinite Cache daemon on a cache server. To install the `edb-icache` daemon on a cache server, deploy the installation wizard on the machine hosting the cache; when the Setup: Select Components window opens, de-select all options except Infinite Cache (as shown in Figure 1.4).

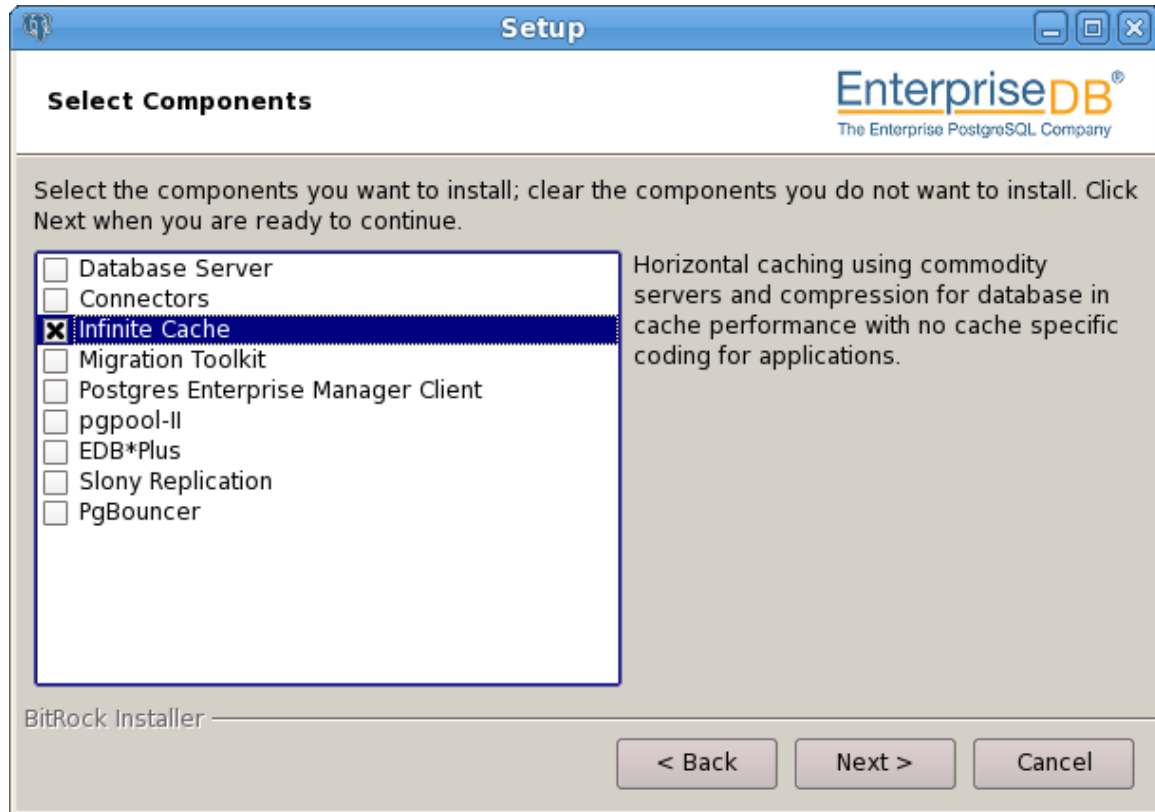


Figure 1.4: Installing only the Infinite Cache Daemon.

The Infinite Cache Daemon option installs the following:

- The `ppas-infinitecache-9.2` service script.
- The Infinite Cache configuration file (`ppas-infinitecache-9.2`).
- A command line tool that allows you to pre-load the cache servers (`edb-icache-warm`).
- The `edb_icache` libraries (code libraries required by the `edb-icache` daemon).

2.2 Configuring the Infinite Cache Server

Configuring Infinite Cache is a three-step process:

- Specify Infinite Cache server settings in the Infinite Cache configuration file.
- Modify the Advanced Server `postgresql.conf` file, enabling Infinite Cache, and specifying connection and compression settings.
- Start the Infinite Cache service.

2.2.1 Modifying Infinite Cache Settings

The Infinite Cache configuration file is named `ppas-infinitecache-9.2`, and contains two parameters and their associated values:

```
PORT=11211
CACHE_SIZE=500
```

To modify a parameter, open the `ppas-infinitecache-9.2` file (located in the `/etc` directory under your Advanced Server installation) with your editor of choice, and modify the parameter values:

`PORT`

Use the `PORT` variable to specify the port where Infinite Cache will listen for connections from Advanced Server.

`CACHE_SIZE`

Use the `CACHE_SIZE` variable to specify the size of the cache (in MB).

2.2.2 Enabling Infinite Cache

The `postgresql.conf` file includes three configuration parameters that control the behavior of Infinite Cache. The `postgresql.conf` file is read each time you start the Advanced Server database server. To modify a parameter, open the `postgresql.conf` file (located in the `$PGDATA` directory) with your editor of choice, and edit the section of the configuration file shown below:

```
# - Infinite Cache
#edb_enable_ichache = off
#edb_ichache_servers = ' ' #'host1:port1,host2,ip3:port3,ip4'
#edb_ichache_compression_level = 6
```

Lines that begin with a pound sign (#) are treated as comments; to enable a given parameter, remove the pound sign and specify a value for the parameter. When you've updated and saved the configuration file, restart the database server for the changes to take effect.

`edb_enable_icache`

Use the `edb_enable_icache` parameter to enable or disable Infinite Cache. When `edb_enable_icache` is set to `on`, Infinite Cache is enabled; if the parameter is set to `off`, Infinite Cache is disabled.

If you set `edb_enable_icache` to `on`, you must also specify a list of cache servers by setting the `edb_icache_servers` parameter (described in the next section).

The default value of `edb_enable_icache` is `off`.

`edb_icache_servers`

The `edb_icache_servers` parameter specifies a list of one or more servers with active `edb-icache` daemons. `edb_icache_servers` is a string value that takes the form of a comma-separated list of *hostname:port* pairs. You can specify each pair in any of the following forms:

- *hostname*
- *IP-address*
- *hostname:portnumber*
- *IP-address:portnumber*

If you do not specify a port number, Infinite Cache assumes that the cache server is listening at port 11211. This configuration parameter will take effect only if `edb_enable_icache` is set to `on`. Use the `edb_icache_servers` parameter to specify a maximum of 128 cache nodes.

`edb_icache_compression_level`

The `edb_icache_compression_level` parameter controls the compression level that is applied to each page before storing it in the distributed Infinite Cache. This parameter must be an integer in the range 0 to 9.

- A compression level of 0 disables compression; it uses no CPU time for compression, but requires more storage space and network resources to process.
- A compression level of 9 invokes the maximum amount of compression; it increases the load on the CPU, but less data flows across the network, so network demand is reduced. Each page takes less room on the Infinite Cache server, so memory requirements are reduced.

- A compression level of 5 or 6 is a reasonable compromise between the amount of compression received and the amount of CPU time invested.

By default, `edb_icache_compression_level` is set to 6.

When Advanced Server reads data from disk, it typically reads the data in 8K increments. If `edb_icache_compression_level` is set to 0, each time Advanced Server sends an 8K page to the Infinite Cache server that page is stored (uncompressed) in 8K of cache memory. If the `edb_icache_compression_level` parameter is set to 9, Advanced Server applies the maximum compression possible before sending it to the Infinite Cache server, so a page that previously took 8K of cached memory might take 2K of cached memory. Exact compression numbers are difficult to predict, as they are dependent on the nature of the data on each page.

The compression level must be set by the superuser and can be changed for the current session while the server is running. The following command disables the compression mechanism for the currently active session:

```
SET edb_icache_compression_level = 0
```

The following example shows a typical collection of Infinite Cache settings:

```
edb_enable_icache           = on
edb_icache_servers          = 'localhost,192.168.2.1:11200,192.168.2.2'
edb_icache_compression_level = 6
```

Please Note: Infinite Cache and the `effective_io_concurrency` parameter can potentially interfere with each other. You should disable asynchronous I/O requests (by setting the value of `effective_io_concurrency` to 0 in the `postgresql.conf` file) if you enable the Infinite Cache feature. By default, `effective_io_concurrency` is set to 1.

2.2.3 Controlling the Infinite Cache Server

Linux

On Linux, the Infinite Cache service script is named `ppas-infinitecache-9.2`. The service script resides in the `/etc/init.d` directory. You can control the Infinite Cache service, or check the status of the service with the following command:

```
/etc/init.d/ppas-infinitecache-9.2 action
```

Where *action* specifies:

- `start` to start the service.
- `stop` to stop the service
- `restart` to stop and then start the service.
- `status` to return the status of the service.

HP-UX

On HP-UX, the Infinite Cache service script is named `ppas-infinitecache-9.2`. The service script resides in the `/sbin/init.d` directory. You can control the Infinite Cache service, or check the status of the service with the following command:

```
/sbin/init.d/ppas-infinitecache-9.2 action
```

Where *action* specifies:

- `start` to start the service.
- `stop` to stop the service
- `restart` to stop and then start the service.
- `status` to return the status of the service.

Solaris

On Solaris, the Infinite Cache service is named `ppas-infinitecache-9_2`, and resides in the `/lib/svc/method` directory. After specifying configuration options, you must manually register and start the Infinite Cache service.

On Solaris 10, enter:

```
svccfg -v import /var/svc/manifest/application/ppas-infinitecache-9_2.xml
```

On Solaris 11, enter:

```
svccfg -v import installation_dir/installer/infinitecache/ppas-infinitecache-9_2.xml
```

After registering and starting the Infinite Cache service, you can use the following command to check the status of the service

```
svcs ppas-infinitecache-9_2
```

You can control the Infinite Cache service with the following command:

```
svcadm action ppas-infinitecache-9_2
```

Where *action* specifies:

- `enable` to start the service.
- `disable` to stop the service
- `restart` to stop and then start the service.

2.3 Dynamically Modifying Infinite Cache Server Nodes

You can dynamically modify the Infinite Cache server nodes; to change the Infinite Cache server configuration, use the `edb_ichache_servers` parameter in the `postgresql.conf` file to:

- specify additional cache information to add a server/s.
- delete server information to remove a server/s.
- specify additional server information and delete existing server information to both add and delete servers during the same reload operation.

After updating the `edb_ichache_servers` parameter in the `postgresql.conf` file, you must reload the configuration parameters for the changes to take effect.

To reload the configuration parameters, navigate through the `Postgres Plus Advanced Server 9.2` menu to the `Expert Configuration` menu, and select the `Reload Configuration` option. If prompted, enter your password to reload the configuration parameters.

Alternatively, you can use the `pg_ctl reload` command to update the server's configuration parameters at the command line:

```
pg_ctl reload -D data_directory
```

Where *data_directory* specifies the complete path to the Advanced Server data directory.

Please Note: If Advanced Server detects a problem with the value specified for the `edb_ichache_servers` parameter during a server reload, it will ignore changes to the parameter and use the last valid parameter value. If you are performing a server restart, and the parameter contains an invalid value, Advanced Server will return an error.

2.4 Controlling the *edb-icache* Daemons

edb-icache is a high-performance memory caching daemon that distributes and stores data in shared buffers. Advanced Server transparently interacts with *edb-icache* daemons to store and retrieve data.

Before starting Advanced Server, the *edb-icache* daemons must be running on each server node. Log into each server and start the *edb-icache* server (on that host) by issuing the following command:

```
# edb-icache -u enterprisedb -d -m 1024
```

Where:

-u

-u specifies the user name

-m

-m specifies the amount of memory to be used by *edb-icache*. The default is 64MB.

-d

-d designates that the service should run in the background

To gracefully kill an *edb-icache* daemon (close any in-use files, flush buffers, and exit), execute the command:

```
# killall -TERM edb-icache
```

If the *edb-icache* daemon refuses to die, you may need to use the following command:

```
# killall -KILL edb-icache
```

2.4.1 Command Line Options

To view the command line options for the `edb-icache` daemon, use the following command from the `edb_Infinite_Cache` subdirectory, located in the Advanced Server installation directory:

```
# edb-icache -h
```

The command line options are:

Parameter	Description
<code>-p <port_number></code>	The TCP port number the Infinite Cache daemon is listening on. The default is 11211.
<code>-U <UDP_number></code>	The UDP port number the Infinite Cache daemon is listening on. The default is 0 (off).
<code>-s <pathname></code>	The Unix socket pathname the Infinite Cache daemon is listening on. If included, the server limits access to the host on which the Infinite Cache daemon is running, and disables network support for Infinite Cache.
<code>-a <mask></code>	The access mask for the Unix socket, in octal form. The default value is 0700.
<code>-l <ip_addr></code>	Specifies the IP address that the daemon is listening on. If an individual address is not specified, the default value is <code>INDRR_ANY</code> ; all IP addresses assigned to the resource are available to the daemon.
<code>-d</code>	Run as a daemon.
<code>-r</code>	Maximize core file limit.
<code>-u <username></code>	Assume the identity of the specified user (when run as root).
<code>-m <numeric></code>	Max memory to use for items in megabytes. Default is 64 MB.
<code>-M</code>	Return error on memory exhausted (rather than removing items).
<code>-c <numeric></code>	Max simultaneous connections. Default is 1024.
<code>-k</code>	Lock down all paged memory. Note that there is a limit on how much memory you may lock. Trying to allocate more than that would fail, so be sure you set the limit correctly for the user you started the daemon with (not for <code>-u <username></code> user; under <code>sh</code> this is done with <code>'ulimit -S -l NUM_KB'</code>).
<code>-v</code>	Verbose (print errors/warnings while in event loop).
<code>-vv</code>	Very verbose (include client commands and responses).
<code>-vvv</code>	Extremely verbose (also print internal state transitions).
<code>-h</code>	Print the help text and exit.
<code>-i</code>	Print memcached and libevent licenses.
<code>-P <file></code>	Save PID in <code><file></code> , only used with <code>-d</code> option.
<code>-f <factor></code>	Chunk size growth factor. Default value is 1.25.
<code>-n <bytes></code>	Minimum space allocated for key+value+flags. Default is 48.
<code>-L</code>	Use large memory pages (if available). Increasing the memory page size could reduce the number of transition look-aside buffer misses and improve the performance. To get large pages from the OS, Infinite Cache will allocate the total item-cache in one large chunk.
<code>-D <char></code>	Use <code><char></code> as the delimiter between key prefixes and IDs. This is used for per-prefix stats reporting. The default is ":" (colon). If this option is specified, stats collection is enabled automatically; if not, then it may be enabled by sending the <code>stats detail on</code> command to the server.
<code>-t <num></code>	Specifies the number of threads to use. Default is 4.
<code>-R</code>	Maximum number of requests per event; this parameter limits the number of requests process for a given connection to prevent starvation, default is 20.
<code>-C</code>	Disable use of CAS (check and set).

-b	Specifies the backlog queue limit, default is 1024.
-B	Specifies the binding protocol. Possible values are ascii, binary or auto; default value is auto.
-I	Override the size of each slab page. Specifies the max item size; default 1 MB, minimum size is 1 k, maximum is 128 MB).

2.4.2 edb-icache-tool

`edb-icache-tool` provides a command line interface that queries the `edb-icache` daemon to retrieve statistical information about a specific cache node. The syntax is:

```
edb-icache-tool <host[:port]> stats
```

host specifies the address of the host that you are querying.

port specifies the port that the daemon is listening on.

`edb-icache-tool` retrieves the statistics described in the following table:

Statistic	Description
accepting_conns	Will this server accept new connection(s)? 1 if yes, otherwise 0.
auth_cmds	Number of authentication commands handled by this server, success or failure.
auth_errors	Number of failed authentications.
bytes	Total number of bytes in use.
bytes_read	Total number of bytes received by this server (from the network).
bytes_written	Total number of bytes sent by this server (to the network).
cas_badval	Number of keys that have been compared and swapped by this server but the comparison (original) value did not match the supplied value.
cas_hits	Number of keys that have been compared and swapped by this server and found present.
cas_misses	Number of keys that have been compared and swapped by this server and not found.
cmd_flush	Cumulative number of flush requests sent to this server.
cmd_get	Cumulative number of read requests sent to this server.
cmd_set	Cumulative number of write requests sent to this server.
conn_yields	Number of times any connection yielded to another due to hitting the <code>edb-icache -R</code> limit.
connection_structures	Number of connection structures allocated by the server.
curr_connections	Number of open connections.
curr_items	Number of items currently stored by the server.
decr_hits	Number of decrement requests satisfied by this server.
decr_misses	Number of decrement requests not satisfied by this server.
delete_hits	Number of delete requests satisfied by this server.
delete_misses	Number of delete requests not satisfied by this server.
evictions	Number of valid items removed from cache to free memory for new items.
get_hits	Number of read requests satisfied by this server.
get_misses	Number of read requests not satisfied by this server.
incr_hits	Number of increment requests satisfied by this server.
incr_misses	Number of increment requests not satisfied by this server.

limit_maxbytes	Number of bytes allocated on this server for storage.
listen_disabled_num	Cumulative number of times this server has hit its connection limit.
pid	Process ID (on cache server).
pointer_size	Default pointer size on host OS (usually 32 or 64).
reclaimed	Number of times an entry was stored using memory from an expired entry.
rusage_user	Accumulated user time for this process (seconds.microseconds).
rusage_system	Accumulated system time for this process (seconds.microseconds).
threads	Number of worker threads requested.
total_time	Number of seconds since this server's base date (usually midnight, January 1, 1970, UTC).
total_connections	Total number of connections opened since the server started running.
total_items	Total number of items stored by this server (cumulative).
uptime	Amount of time that server has been active.
version	edb-icache version.

In the following example, edb-icache-tool retrieves statistical information about an Infinite Cache server located at the address, 192.168.23.85 and listening on port 11213:

```
# edb-icache-tool 192.168.23.85:11213 stats
```

```
Field                Value
accepting_conns      1
auth_cmds             0
auth_errors          0
bytes                52901223
bytes_read            188383848
bytes_written         60510385
cas_badval           0
cas_hits              0
cas_misses            0
cmd_flush             1
cmd_get               53139
cmd_set               229120
conn_yields           0
connection_structures 34
curr_connections      13
curr_items            54953
decr_hits             0
decr_misses           0
delete_hits           0
delete_misses         0
evictions             0
get_hits              52784
get_misses            355
incr_hits             0
incr_misses           0
limit_maxbytes        314572800
listen_disabled_num   0
pid                   7226
pointer_size          32
reclaimed             0
rusage_system         10.676667
rusage_user           3.068191
threads               4
time                  1320919080
total_connections     111
```

```
total_items      229120
uptime           7649
version          1.4.5
```

2.5 Warming the edb-icache Servers

When Advanced Server starts, the primary and secondary caches are empty. When Advanced Server processes a client request, Advanced Server reads the required data from disk and stores a copy in each cache. You can improve server performance by *warming* (or pre-loading) the data into the memory cache before a client asks for it.

There are two advantages to warming the cache. Advanced Server will find data in the cache the first time it is requested by a client application, instead of waiting for it to be read from disk. Also, manually warming the cache with the data that your applications are most likely to need saves time by avoiding future random disk reads. If you don't warm the cache at startup, Postgres Plus Advanced Server performance may not reach full speed until the client applications happen to load commonly used data into the cache.

There are several ways to load pages to warm the Infinite Cache server nodes. You can:

- Use the `edb_icache_warm` utility to warm the caches from the command line.
- Use the `edb_icache_warm()` function from within `edb-psql`.
- Use the `edb_icache_warm()` function via scripts to warm the cache.

While it is not necessary to re-warm the cache after making changes to an existing cache configuration, re-warming the cache can improve performance by bringing the new configuration of cache servers up-to-date.

2.5.1 The `edb_icache_warm()` Function

The `edb_icache_warm()` function comes in two variations; the first variation warms not only the table, but any indexes associated with the table. If you use the second variation, you must make additional calls to warm any associated indexes.

The first form of the `edb_icache_warm()` function warms the given table and any associated indexes into the cache. The signature is:

```
edb_icache_warm(table_name)
```

You may specify `table_name` as a table name, OID, or `regclass` value.

```
# edb-psql edb -c "select edb_icache_warm('accounts')"
```

When you call the first form of `edb_icache_warm()`, Advanced Server reads each page in the given table, compresses the page (if configured to do so), and then sends the compressed data to an Infinite Cache server. `edb_icache_warm()` also reads, compresses, and caches each page in each index defined for the given table.

The second form of the `edb_icache_warm()` function warms the pages that contain the specified range of bytes into the cache. The signature of the second form is:

```
edb_icache_warm(table-spec, startbyte, endbyte):
```

You must make subsequent calls to specify indexes separately when using this form of the `edb_icache_warm()` function.

```
# edb-psql edb -c "select edb_icache_warm('accounts', 1, 10000)"
```

The `edb_icache_warm()` function is typically called by a utility program (such as the `edb_icache_warm` utility) to spread the warming process among multiple processes that operate in parallel.

2.5.2 Using the `edb_icache_warm` Utility

You can use the `edb_icache_warm` command-line utility to load the cache servers with specified tables, allowing fast access to relevant data from the cache.

The syntax for `edb_icache_warm` is:

```
# edb_icache_warm -d database -t tablename
```

The only required parameter is *tablename*. *tablename* can be specified with or without the `-t` option. All other parameters are optional; if omitted, default values are inferred from Advanced Server environment variables.

The options for `edb_icache_warm` are:

Option	Variable	Description
<code>-h</code>	<i>hostname</i>	The name of the host running Advanced Server. Include this parameter if you are running Advanced Server on a remote host. The default value is PGHOST.
<code>-p</code>	<i>portname</i>	Port in use by Advanced Server. Default value is PGPORT.
<code>-j</code>	<i>process count</i>	Number of (parallel) processes used to warm the cache. The default value is 1.
<code>-U</code>	<i>username</i>	The Advanced Server username. Unless specified, this defaults to PGUSER.
<code>-d</code>	<i>database</i>	The name of database containing the tables to be warmed. Default value is PGDATABASE.
<code>-t</code>	<i>tablename</i>	Name of table to be warmed. The index for the table is also warmed. Required.

2.6 Retrieving Statistics from Infinite Cache

2.6.1 Using edb_icache_stats()

You can view Infinite Cache statistics by using the `edb_icache_stats()` function at the `edb-psql` command line (or any other query tool).

The `edb_icache_stats()` function returns a result set that reflects the state of an Infinite Cache node or nodes and the related usage statistics. The result set includes:

Statistic	Description
<i>hostname</i>	Host name (or IP address) of server
<i>port</i>	Port number at which edb-icache daemon is listening
<i>state</i>	Health of this server
<i>write_failures</i>	Number of write failures
<i>bytes</i>	Total number of bytes in use
<i>bytes_read</i>	Total number of bytes received by this server (from the network)
<i>bytes_written</i>	Total number of bytes sent by this server (to the network)
<i>cmd_get</i>	Cumulative number of read requests sent to this server
<i>cmd_set</i>	Cumulative number of write requests sent to this server
<i>connection_structures</i>	Number of connection structures allocated by the server
<i>curr_connections</i>	Number of open connections
<i>curr_items</i>	Number of items currently stored by the server
<i>evictions</i>	Number of valid items removed from cache to free memory for new items
<i>get_hits</i>	Number of read requests satisfied by this server
<i>get_misses</i>	Number of read requests not satisfied by this server
<i>limit_maxbytes</i>	Number of bytes allocated on this server for storage
<i>pid</i>	Process ID (on cache server)
<i>pointer_size</i>	Default pointer size on host OS (usually 32 or 64)
<i>rusage_user</i>	Accumulated user time for this process (seconds.microseconds)
<i>rusage_system</i>	Accumulated system time for this process (seconds.microseconds)
<i>threads</i>	Number of worker threads requested
<i>total_time</i>	Number of seconds since this server's base date (usually midnight, January 1, 1970, UTC)
<i>total_connections</i>	Total number of connections opened since the server started running
<i>total_items</i>	Total number of items stored by this server (cumulative)
<i>uptime</i>	Amount of time that server has been active
<i>version</i>	edb-icache version

You can use SQL queries to view Infinite Cache statistics. To view the server status of all Infinite Cache nodes:

```
SELECT hostname, port, state FROM edb_icache_stats()
```

```

hostname      | port  | state
-----+-----+-----
192.168.23.85 | 11211 | UNHEALTHY
192.168.23.85 | 11212 | ACTIVE
(2 rows)
```

Use the following command to view complete statistics (shown here using `edb-psql`'s expanded display mode, `\x`) for a specified node:

```
SELECT * FROM edb_icache_stats() WHERE hostname = '192.168.23.85:11211'
```

```

-[RECORD 1]-----+-----
hostname           | 192.168.23.85
port               | 11211
state              | ACTIVE
write_failures     | 0
bytes              | 225029460
bytes_read         | 225728252
bytes_written      | 192806774
cmd_get            | 23313
cmd_set            | 27088
connection_structures | 53
curr_connections   | 3
curr_items         | 27088
evictions          | 0
get_hits           | 23266
get_misses         | 47
limit_maxbytes     | 805306368
pid                | 4240
pointer_size       | 32
rusage_user        | 0.481926
rusage_system      | 1.583759
threads            | 1
total_time         | 1242199782
total_connections  | 66
total_items        | 27088
uptime             | 714
version            | 1.2.6

```

2.6.2 edb_icache_server_list

The `edb_icache_server_list` view exposes information about the status and health of all Infinite Cache servers listed in the `edb_icache_servers` GUC. The `edb_icache_server_list` view is created using the `edb_icache_stats()` API. The view exposes the following information for each server:

Statistic	Description
<i>hostname</i>	Host name (or IP address) of server
<i>port</i>	Port number at which edb-icache daemon is listening
<i>state</i>	Health of this server
<i>write_failures</i>	Number of write failures
<i>total_memory</i>	Number of bytes allocated to the cache on this server
<i>memory_used</i>	Number of bytes currently used by the cache
<i>memory_free</i>	Number of unused bytes remaining in the cache
<i>hit_ratio</i>	Percentage of cache hits

The `state` column will contain one of the following four values, reflecting the health of the given server:

Server State	Description
Active	The server is known to be up and running.
Unhealthy	An error occurred while interacting with the cache server. Postgres will attempt to re-establish the connection with the server.
Offline	Postgres can no longer contact the given server.
Manual Offline	You have taken the server offline with the <code>edb_icache_server_enable()</code> function.

Use the following `SELECT` statement to return the health of each node in the Infinite Cache server farm:

```
SELECT hostname, port, state FROM edb_icache_server_list
```

```

  hostname | port | state
-----+-----+-----
 192.168.23.85 | 11211 | ACTIVE
 192.168.23.85 | 11212 | ACTIVE
(2 rows)
```

Use the following command to view complete details about a specific Infinite Cache node (shown here using `edb-psql`'s `\x` expanded-view option):

```
SELECT * FROM edb_icache_server_list WHERE hostname = '192.168.23.85:11211'
```

```

-[RECORD 1]-----+-----
hostname          | 192.168.23.85
port              | 11211
state             | ACTIVE
write_failures    | 0
total_memory      | 805306368
memory_used       | 225029460
memory_free       | 580276908
hit_ratio         | 99.79
```

2.7 Retrieving Table Statistics

Advanced Server provides six system views that contain statistical information on a per-table basis. The views are:

- `pg_statio_all_tables`
- `pg_statio_sys_tables`
- `pg_statio_user_tables`
- `pg_statio_all_indexes`
- `pg_statio_sys_indexes`
- `pg_statio_user_indexes`

You can use standard SQL queries to view and compare the information stored in the views. The views contain information that will allow you to observe the effectiveness of the Advanced Server buffer cache and the icache servers.

2.7.1 pg_statio_all_tables

The `pg_statio_all_tables` view contains one row for each table in the database. The view contains the following information:

Column Name	Description
<code>relid</code>	The OID of the table.
<code>schemaname</code>	The name of the schema that the table resides in.
<code>relname</code>	The name of the table.
<code>heap_blks_read</code>	The number of heap blocks read.
<code>heap_blks_hit</code>	The number of heap blocks hit.
<code>heap_blks_icache_hit</code>	The number of heap blocks found on an icache server.
<code>idx_blks_read</code>	The number of index blocks read.
<code>idx_blks_hit</code>	The number of index blocks hit.
<code>idx_blks_icache_hit</code>	The number of index blocks found on an icache server.
<code>toast_blks_read</code>	The number of toast blocks read.
<code>toast_blks_hit</code>	The number of toast blocks hit.
<code>toast_blks_icache_hit</code>	The number of toast blocks found on an icache server.
<code>tidx_blks_read</code>	The number of index toast blocks read.
<code>tidx_blks_hit</code>	The number of index toast blocks hit.
<code>tidx_blks_icache_hit</code>	The number of index toast blocks found on an icache server.

You can execute a simple query to view performance statistics for a specific table:

```
SELECT * FROM pg_statio_all_tables WHERE relname='jobhist';

-[ RECORD 1 ]-----+-----
relid          | 16402
schemaname     | public
relname        | jobhist
heap_blks_read | 1
heap_blks_hit  | 51
heap_blks_icache_hit | 0
idx_blks_read  | 2
idx_blks_hit   | 17
idx_blks_icache_hit | 0
toast_blks_read |
toast_blks_hit |
toast_blks_icache_hit |
tidx_blks_read |
tidx_blks_hit  |
tidx_blks_icache_hit |
```

Or, you can view the statistics by activity level. The following example displays the statistics for the ten tables that have the greatest `heap_blks_icache_hit` activity:

```
SELECT * FROM pg_statio_all_tables ORDER BY heap_blks_icache_hit DESC LIMIT 10;

relid      schemaname      relname
heap_blks_read  heap_blks_hit  heap_blks_icache_hit
idx_blks_read   idx_blks_hit   idx_blks_icache_hit
toast_blks_read toast_blks_hit toast_blks_icache_hit
tidx_blks_read  tidx_blks_hit  tidx_blks_icache_hit
-----
```

```

16390      public      pgbench_accounts
    264105          71150      81498
    13171          282541     18053

1259      pg_catalog  pg_class
    22          2904      18
    14          3449      11

1249      pg_catalog  pg_attribute
    49          1619      16
    17          2841      13

1255      pg_catalog  pg_proc
    38          276       11
    33          682       16
    0           0         0
    0           0         0

2619      pg_catalog  pg_statistic
    20          295       8
    4           436       4
    0           0         0
    0           0         0

2617      pg_catalog  pg_operator
    20          293       8
    19          791      10

2602      pg_catalog  pg_amop
    10          721       6
    13          1154      13

2610      pg_catalog  pg_index
    10          633       6
    8           719       8

1247      pg_catalog  pg_type
    17          235       5
    12          433       4

2615      pg_catalog  pg_namespace
    4           260       4
    6           330       4
    0           0         0
    0           0         0
(10 rows)

```

2.7.2 pg_statio_sys_tables

The `pg_statio_sys_tables` view contains one row for each table in a system-defined schema. The statistical information included in this view is the same as for `pg_statio_all_tables`.

2.7.3 pg_statio_user_tables

The `pg_statio_user_tables` view contains one row for each table in a user-defined schema. The statistical information in this view is the same as for `pg_statio_all_tables`.

2.7.4 pg_statio_all_indexes

The `pg_statio_all_indexes` view contains one row for each index in the current database. The view contains the following information:

Column Name	Description
<code>relid</code>	The OID of the indexed table
<code>indexrelid</code>	The OID of the index.
<code>schemaname</code>	The name of the schema that the table resides in.
<code>relname</code>	The name of the table.
<code>indexrelname</code>	The name of the index
<code>idx_blks_read</code>	The number of index blocks read.
<code>idx_blks_hit</code>	The number of index blocks hit.
<code>idx_blks_ichache_hit</code>	The number of index blocks found on an icache server.

You can execute a simple query to view performance statistics for the indexes on a specific table:

```
SELECT * FROM pg_statio_all_indexes WHERE relname='pg_attribute';

-[ RECORD 1 ]-----+-----
relid          | 1249
indexrelid     | 2658
schemaname     | pg_catalog
relname        | pg_attribute
indexrelname    | pg_attribute_relid_attnam_index
idx_blks_read  | 10
idx_blks_hit   | 1200
idx_blks_ichache_hit | 0
-[ RECORD 2 ]-----+-----
relid          | 1249
indexrelid     | 2659
schemaname     | pg_catalog
relname        | pg_attribute
indexrelname    | pg_attribute_relid_attnum_index
idx_blks_read  | 12
idx_blks_hit   | 3917
idx_blks_ichache_hit | 0
```

The result set from the query includes the statistical information for two indexes; the `pg_attribute` table has two indexes.

You can also view the statistics by activity level. The following example displays the statistics for the ten indexes that have the greatest `idx_blks_ichache_hit` activity:

```
SELECT * FROM pg_statio_all_indexes ORDER BY idx_blks_icache_hit DESC LIMIT 10;
```

relid	indexrelid	schemaname	relname	idx_blks_read	idx_blks_hit	idx_blks_icache_hit
indexrelname						
16390	16401	public	pgbench_accounts	13171	282541	18053
pgbench_accounts_pkey						
1249	2659	pg_catalog	pg_attribute	14	2749	13
pg_attr_relid_attnum_index						
1255	2690	pg_catalog	proc	16	580	12
pg_proc_oid_index						
1259	2663	pg_catalog	pg_class	10	2019	7
pg_class_relname_nsp_index						
2602	2654	pg_catalog	pg_amop	7	453	7
pg_amop_opr_fam_index						
2603	2655	pg_catalog	pg_amproc	6	605	6
pg_amproc_fam_proc_index						
2617	2688	pg_catalog	pg_operator	7	452	6
pg_operator_oid_index						
2602	2653	pg_catalog	pg_amop	6	701	6
pg_amop_fam_strat_index						
2615	2684	pg_catalog	pg_namespace	4	328	4
pg_namespace_nspname_index						
1262	2672	pg_catalog	pg_database	4	254	4
pg_database_oid_index						

2.7.5 pg_statio_sys_indexes

The `pg_statio_sys_indexes` view contains one row for each index on the system tables. The statistical information in this view is the same as in `pg_statio_all_indexes`.

2.7.6 pg_statio_user_indexes

The `pg_statio_user_indexes` view contains one row for each index on a table that resides in a user-defined schema. The statistical information in this view is the same as in `pg_statio_all_indexes`.

2.8 edb_icache_server_enable()

You can use the `edb_icache_server_enable()` function to take the Infinite Cache server offline for maintenance or other planned downtime. The syntax is:

```
void edb_icache_server_enable(host TEXT, port INTEGER, online BOOL)
```

host specifies the host that you want to disable. The host name may be specified by name or numeric address.

port specifies the port number that the Infinite Cache server is listening on.

online specifies the state of the Infinite Cache server. The value of *online* must be `true` or `false`.

To take a server offline, specify the host that you want to disable, the port number that the Infinite Cache server is listening on, and `false`. To bring the Infinite Cache server back online, specify the host name and port number, and pass a value of `true`.

The state of a server taken offline with the `edb_icache_server_enable()` function is `MANUAL OFFLINE`. Postgres Plus Advanced Server will not automatically reconnect to an Infinite Cache server that you have taken offline with

`edb_icache_server_enable(..., false)`; you must bring the server back online by calling `edb_icache_server_enable(..., true)`.

2.9 Infinite Cache Log Entries

When you start Advanced Server, a message that includes Infinite Cache status, cache node count and cache nodes size is written to the server log. The following example shows the server log for an active Infinite Cache installation with two 750 MB cache servers:

```
** EnterpriseDB Dynamic Tuning Agent*****
*      System Utilization: 66 %                      *
*      Autovacuum Naptime: 60      Seconds            *
*      Infinite Cache: on                             *
*      Infinite Cache Servers: 2                      *
*      Infinite Cache Size: 1.500  GB                 *
*****
```


2.10 Allocating Memory to the Cache Servers

As mentioned earlier in this document, each computer imposes a limit on the amount of *physical* memory that you can install. However, modern operating systems typically simulate a larger *address space* so that programs can transparently access more memory than is actually installed. This "virtual memory" allows a computer to run multiple programs which may simultaneously require more memory than is physically available. For example, you may run an e-mail client, a web browser, and a database server which each require 1GB of memory on a machine that contains only 2GB of physical RAM. When the operating system runs out of physical memory, it starts swapping bits and pieces of the currently running programs to disk to make room to satisfy your current demand for memory.

This can bring your system to a grinding halt.

Since the primary goal of Infinite Cache is to improve performance by limiting disk I/O, you should avoid dedicating so much memory to Infinite Cache that the operating system must start swapping data to disk. If the operating system begins to swap to disk, you lose the benefits offered by Infinite Cache.

The overall demand for physical memory can vary throughout the day; if the server is frequently idle, you may never encounter swapping. If you have dedicated a large portion of physical memory to the cache, and system usage increases, the operating system may start swapping. To get the best performance and avoid disk swapping, dedicate a server node to Infinite Cache so other applications on that computer will not compete for physical memory.

3 Dynatune

Postgres Plus Advanced Server supports dynamic tuning of the database server to make the optimal usage of the system resources available on the host machine on which it is installed. The two parameters that control this functionality are located in the `postgresql.conf` file. These parameters are:

- `edb_dynatune`
- `edb_dynatune_profile`

3.1.1 edb_dynatune

`edb_dynatune` determines how much of the host system's resources are to be used by the database server based upon the host machine's total available resources and the intended usage of the host machine.

When Postgres Plus Advanced Server is initially installed, the `edb_dynatune` parameter is set in accordance with the selected usage of the host machine on which it was installed - i.e., development machine, mixed use machine, or dedicated server. For most purposes, there is no need for the database administrator to adjust the various configuration parameters in the `postgresql.conf` file in order to improve performance.

You can change the value of the `edb_dynatune` parameter after the initial installation of Postgres Plus Advanced Server by editing the `postgresql.conf` file. The postmaster must be restarted in order for the new configuration to take effect.

The `edb_dynatune` parameter can be set to any integer value between 0 and 100, inclusive. A value of 0, turns off the dynamic tuning feature thereby leaving the database server resource usage totally under the control of the other configuration parameters in the `postgresql.conf` file.

A low non-zero, value (e.g., 1 - 33) dedicates the least amount of the host machine's resources to the database server. This setting would be used for a development machine where many other applications are being used.

A value in the range of 34 - 66 dedicates a moderate amount of resources to the database server. This setting might be used for a dedicated application server that may have a fixed number of other applications running on the same machine as Postgres Plus Advanced Server.

The highest values (e.g., 67 - 100) dedicate most of the server's resources to the database server. This setting would be used for a host machine that is totally dedicated to running Postgres Plus Advanced Server.

Once a value of `edb_dynatune` is selected, database server performance can be further fine-tuned by adjusting the other configuration parameters in the `postgresql.conf` file. Any adjusted setting overrides the corresponding value chosen by `edb_dynatune`. You can change the value of a parameter by un-commenting the configuration parameter, specifying the desired value, and restarting the database server.

3.1.2 edb_dynatune_profile

The `edb_dynatune_profile` parameter is used to control tuning aspects based upon the expected workload profile on the database server. This parameter takes effect upon startup of the database server.

The possible values for `edb_dynatune_profile` are:

Value	Usage
<code>oltp</code>	Recommended when the database server is processing heavy online transaction processing workloads.
<code>reporting</code>	Recommended for database servers used for heavy data reporting.
<code>mixed</code>	Recommended for servers that provide a mix of transaction processing and data reporting.

4 Dynamic Runtime Instrumentation Tools Architecture (DRITA)

Note: *This information is also included in the Oracle® Compatibility Developer's Guide.*

The Dynamic Runtime Instrumentation Tools Architecture (DRITA) allows a DBA to query catalog views to determine the *wait events* that affect the performance of individual sessions or the system as a whole. DRITA records the number of times each event occurs as well as the time spent waiting; you can use this information to diagnose performance problems.

DRITA compares *snapshots* to evaluate the performance of a system. A snapshot is a saved set of system performance data at a given point in time. Each snapshot is identified by a unique ID number; you can use snapshot ID numbers with DRITA reporting functions to return system performance statistics.

DRITA consumes minimal system resources.

4.1.1 Initialization Parameters

DRITA includes a configuration parameter, `timed_statistics`, to control the collection of timing data. This is a dynamic parameter that can be set in the `postgresql.conf` file or while a session is in progress. The valid values are `TRUE` or `FALSE`; the default value is `FALSE`.

4.1.2 Setting up and Using DRITA

First, take a beginning snapshot. The beginning snapshot will be compared to a later snapshot to gauge system performance. To take a beginning snapshot:

```
SELECT * from edbsnap()
```

Then, run the workload that you would like to evaluate; when the workload has completed (or at a strategic point during the workload), take an ending snapshot:

```
SELECT * from edbsnap()
```

4.2 DRITA Functions

4.2.1 get_snaps()

The `get_snaps()` function returns a list of the current snapshots. The signature is:

```
get_snaps()
```

The following example demonstrates using the `get_snaps()` function to display a list of snapshots:

```
edb=# SELECT * FROM get_snaps();
      get_snaps
-----
 1  11-FEB-10 10:41:05.668852
 2  11-FEB-10 10:42:27.26154
 3  11-FEB-10 10:45:48.999992
 4  11-FEB-10 11:01:58.345163
 5  11-FEB-10 11:05:14.092683
 6  11-FEB-10 11:06:33.151002
 7  11-FEB-10 11:11:16.405664
 8  11-FEB-10 11:13:29.458405
 9  11-FEB-10 11:23:57.595916
10  11-FEB-10 11:29:02.214014
11  11-FEB-10 11:31:44.244038
(11 rows)
```

The first column in the list of snapshots contains the session identifier; the DRITA functions use the session identifier to operate on a specific snapshot.

4.2.2 sys_rpt()

The `sys_rpt()` function returns system wait information. The signature is:

```
sys_rpt(beginning_id, ending_id, top_n)
```

Parameters

beginning_id

beginning_id is an integer value that represents the beginning session identifier.

ending_id

ending_id is an integer value that represents the ending session identifier.

top_n

top_n represents the number of rows to return

This example demonstrates a call to the `sys_rpt()` function:

```
edb=# SELECT * FROM sys_rpt(9, 10, 10);
```

sys_rpt			
WAIT NAME	COUNT	WAIT TIME	% WAIT
wal write	21250	104.723772	36.31
db file read	121407	72.143274	25.01
wal flush	84185	51.652495	17.91
wal file sync	712	29.482206	10.22
infinitecache write	84178	15.814444	5.48
db file write	84177	14.447718	5.01
infinitecache read	672	0.098691	0.03
db file extend	190	0.040386	0.01
query plan	52	0.024400	0.01
wal insert lock acquire	4	0.000837	0.00
(12 rows)			

4.2.3 sess_rpt()

The `sess_rpt()` function returns session wait information. The signature is:

```
sess_rpt(beginning_id, ending_id, top_n)
```

Parameters

beginning_id

beginning_id is an integer value that represents the beginning session identifier.

ending_id

ending_id is an integer value that represents the ending session identifier.

top_n

top_n represents the number of rows to return

The following example demonstrates a call to the `sess_rpt()` function:

sess_rpt							
ID	USER	WAIT NAME	COUNT	TIME (ms)	%WAIT SES	%WAIT ALL	
17373	enterprise	db file read	30	0.175713	85.24	85.24	
17373	enterprise	query plan	18	0.014930	7.24	7.24	
17373	enterprise	wal flush	6	0.004067	1.97	1.97	
17373	enterprise	wal write	1	0.004063	1.97	1.97	
17373	enterprise	wal file sync	1	0.003664	1.78	1.78	
17373	enterprise	infinitecache read	38	0.003076	1.49	1.49	

```

17373 enterprise infinitecache write      5      0.000548    0.27      0.27
17373 enterprise db file extend          190    0.04.386    0.03      0.03
17373 enterprise db file write           5      0.000082    0.04      0.04
17373 enterprise wal write lock acquire  0      0.000000    0.00      0.00
17373 enterprise bgwriter comm lock ac   0      0.000000    0.00      0.00
(13 rows)

```

4.2.4 sessid_rpt()

The `sessid_rpt()` function returns session ID information for a specified backend. The signature is:

```
sessid_rpt(beginning_id, ending_id, backend_id)
```

Parameters

beginning_id

beginning_id is an integer value that represents the beginning session identifier.

ending_id

ending_id is an integer value that represents the ending session identifier.

backend_id

backend_id is an integer value that represents the backend identifier.

The following code sample demonstrates a call to `sessid_rpt()`:

```

SELECT * FROM sessid_rpt(18, 19, 17373);

```

sessid_rpt						
ID	USER	WAIT NAME	COUNT	TIME (ms)	%WAIT SES	%WAIT ALL
17373	enterprise	db file read	30	0.175713	85.24	85.24
17373	enterprise	query plan	18	0.014930	7.24	7.24
17373	enterprise	wal flush	6	0.004067	1.97	1.97
17373	enterprise	wal write	1	0.004063	1.97	1.97
17373	enterprise	wal file sync	1	0.003664	1.78	1.78
17373	enterprise	infinitecache read	38	0.003076	1.49	1.49
17373	enterprise	infinitecache write	5	0.000548	0.27	0.27
17373	enterprise	db file extend	190	0.040386	0.03	0.03
17373	enterprise	db file write	5	0.000082	0.04	0.04
17373	enterprise	wal write lock acquire	0	0.000000	0.00	0.00
17373	enterprise	bgwriter comm lock ac	0	0.000000	0.00	0.00

(13 rows)

4.2.5 sesshist_rpt()

The `sesshist_rpt()` function returns session wait information for a specified backend. The signature is:

```
sesshist_rpt(snapshot_id, session_id)
```

Parameters

snapshot_id

snapshot_id is an integer value that identifies the snapshot.

session_id

session_id is an integer value that represents the session.

The following example demonstrates a call to the `sesshist_rpt()` function:

```
edb=# SELECT * FROM sesshist_rpt (9, 5531);
           sesshist_rpt
```

ID	USER	SEQ	WAIT NAME		
ELAPSED(ms)	File	Name	# of Blk	Sum of Blks	
5531	enterprise	1	db file read		
18546	14309	session_waits_pk	1	1	
5531	enterprise	2	infinitemcache read		
125	14309	session_waits_pk	1	1	
5531	enterprise	3	db file read		
376	14304	edb\$session_waits	0	1	
5531	enterprise	4	infinitemcache read		
166	14304	edb\$session_waits	0	1	
5531	enterprise	5	db file read		
7978	1260	pg_authid	0	1	
5531	enterprise	6	infinitemcache read		
154	1260	pg_authid	0	1	
5531	enterprise	7	db file read		
628	14302	system_waits_pk	1	1	
5531	enterprise	8	infinitemcache read		
463	14302	system_waits_pk	1	1	
5531	enterprise	9	db file read		
3446	14297	edb\$system_waits	0	1	
5531	enterprise	10	infinitemcache read		
187	14297	edb\$system_waits	0	1	
5531	enterprise	11	db file read		
14750	14295	snap_pk	1	1	
5531	enterprise	12	infinitemcache read		
416	14295	snap_pk	1	1	
5531	enterprise	13	db file read		
7139	14290	edb\$snap	0	1	
5531	enterprise	14	infinitemcache read		
158	14290	edb\$snap	0	1	
5531	enterprise	15	db file read		
27287	14288	snapshot_num_seq	0	1	
5531	enterprise	16	infinitemcache read		


```

180          14288 snapshot_num_seq      0          1
5531 enterprise 17      query plan
26          0      N/A              0          0
5531 enterprise 18      db file read
84552       16396 pgbench_accounts    4358      1
5531 enterprise 19      infinitecache read
226         16396 pgbench_accounts    4358      1
5531 enterprise 20      db file read
334838      16401 pgbench_accounts_pke 7792      1
5531 enterprise 21      infinitecache read
213         16401 pgbench_accounts_pke 7792      1
5531 enterprise 22      db file read
52619       16396 pgbench_accounts    24829     1
5531 enterprise 23      infinitecache read
210         16396 pgbench_accounts    24829     1
5531 enterprise 24      infinitecache read
216         16401 pgbench_accounts_pke 13460     1
5531 enterprise 25      db file read
13925       16396 pgbench_accounts    27695     1
(27 rows)

```

4.2.6 purgesnap()

The `purgesnap()` function purges a range of snapshots from the snapshot tables. The signature is:

```
purgesnap(beginning_id, ending_id)
```

Parameters

beginning_id

beginning_id is an integer value that represents the beginning session identifier.

ending_id

ending_id is an integer value that represents the ending session identifier.

`purgesnap()` removes all snapshots between *beginning_id* and *ending_id* (inclusive):

```

SELECT * FROM purgesnap(6, 9);

      purgesnap
-----
Snapshots in range 6 to 9 deleted.
(1 row)

```

A call to the `get_snaps()` function after executing the example shows that snapshots 6 through 9 have been purged from the snapshot tables:

```
edb=# SELECT * FROM get_snaps();
      get_snaps
-----
 1  11-FEB-10 10:41:05.668852
 2  11-FEB-10 10:42:27.26154
 3  11-FEB-10 10:45:48.999992
 4  11-FEB-10 11:01:58.345163
 5  11-FEB-10 11:05:14.092683
10  11-FEB-10 11:29:02.214014
11  11-FEB-10 11:31:44.244038
(7 rows)
```

4.2.7 truncsnap()

Use the `truncsnap()` function to delete all records from the snapshot table. The signature is:

```
truncsnap()
```

For example:

```
SELECT * FROM truncsnap();
      truncsnap
-----
Snapshots truncated.
(1 row)
```

A call to the `get_snaps()` function after calling the `truncsnap()` function shows that all records have been removed from the snapshot tables:

```
SELECT * FROM get_snaps();
      get_snaps
-----
(0 rows)
```

4.3 Simulating Statspack AWR Reports

The functions described in this section return information comparable to the information contained in an Oracle Statspack/AWR (Automatic Workload Repository) report. When taking a snapshot, performance data from system catalog tables is saved into history tables. The reporting functions listed below report on the differences between two given snapshots.

- `stat_db_rpt()`
- `stat_tables_rpt()`
- `statio_tables_rpt()`
- `stat_indexes_rpt()`
- `statio_indexes_rpt()`

The reporting functions can be executed individually or you can execute all five functions by calling the `edbreport()` function.

4.3.1 edbreport()

The `edbreport()` function includes data from the other reporting functions, plus additional system information. The signature is:

```
edbreport(beginning_id, ending_id)
```

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

The following code sample demonstrates a call to the `edbreport()` function:

```
edb=# SELECT * FROM edbreport(9, 10);

edbreport
-----
EnterpriseDB Report for database edb          23-AUG-12
Version: EnterpriseDB 9.2.0.0 on i686-pc-linux-gnu, compiled by gcc (GCC)
4.1.2 20080704 (Red Hat 4.1.2-52), 32-bit

Begin snapshot: 9 at 23-AUG-12 13:45:07.165123
End snapshot:   10 at 23-AUG-12 13:45:35.653036
```

Size of database edb is 155 MB		
Tablespace: pg_default Size: 179 MB Owner: enterprisedb		
Tablespace: pg_global Size: 435 kB Owner: enterprisedb		
Schema: pg_toast_temp_1	Size: 0 bytes	Owner:
enterprisedb		
Schema: public	Size: 0 bytes	Owner:
enterprisedb		
Schema: enterprisedb	Size: 143 MB	Owner:
enterprisedb		
Schema: pgagent	Size: 192 kB	Owner:
enterprisedb		
Schema: dbms_job_procedure	Size: 0 bytes	Owner:
enterprisedb		
Top 10 Relations by pages		
TABLE	RELPGAGES	

pgbench_accounts	15874	
pg_proc	102	
edb\$statio_all_indexes	73	
edb\$stat_all_indexes	73	
pg_attribute	67	
pg_depend	58	
edb\$statio_all_tables	49	
edb\$stat_all_tables	47	
pgbench_tellers	37	
pg_description	32	
Top 10 Indexes by pages		
INDEX	RELPGAGES	

pgbench_accounts_pkey	2198	
pg_depend_depender_index	32	
pg_depend_reference_index	31	
pg_proc_proname_args_nsp_index	30	
pg_attribute_relid_attnam_index	23	
pg_attribute_relid_attnum_index	17	
pg_description_o_c_o_index	15	
edb\$statio_idx_pk	11	
edb\$stat_idx_pk	11	
pg_proc_oid_index	9	
Top 10 Relations by DML		
SCHEMA	RELATION	UPDATES
DELETES	INSERTS	

enterprisedb	pgbench_accounts	10400
1000000		0
enterprisedb	pgbench_tellers	10400
100		0
enterprisedb	pgbench_branches	10400
10		0
enterprisedb	pgbench_history	0
10400		0
pgagent	pga_jobclass	0
		0
6		

```

pgagent      pga_exception      0      0
0
pgagent      pga_job      0      0
0
pgagent      pga_jobagent      0      0
0
pgagent      pga_joblog      0      0
0
pgagent      pga_jobstep      0      0
0

DATA from pg_stat_database

DATABASE      NUMBACKENDS      XACT COMMIT      XACT ROLLBACK      BLKS READ      BLKS HIT
BLKS ICACHE HIT      HIT RATIO      ICACHE HIT RATIO
-----
edb      0      142      0      78      10446      0
99.26      0.00

DATA from pg_bufferscache not included because pg_bufferscache is not
installed

DATA from pg_stat_all_tables ordered by seq scan

SCHEMA      RELATION      SEQ SCAN      REL TUP READ
IDX SCAN      IDX TUP READ INS      UPD      DEL
-----
pg_catalog      pg_class      16      7162
546      319      0      1      0
pg_catalog      pg_am      13      13
0      0      0      0      0
pg_catalog      pg_database      4      16
42      42      0      0      0
pg_catalog      pg_index      4      660
145      149      0      0      0
pg_catalog      pg_namespace      4      100
49      49      0      0      0
sys      edb$snap      1      9
0      0      1      0      0
pg_catalog      pg_authid      1      1
25      25      0      0      0
sys      edb$session_wait_history      0      0
0      0      50      0      0
sys      edb$session_waits      0      0
0      0      2      0      0
sys      edb$stat_all_indexes      0      0
0      0      165      0      0

DATA from pg_stat_all_tables ordered by rel tup read

SCHEMA      RELATION      SEQ SCAN      REL TUP READ
IDX SCAN      IDX TUP READ INS      UPD      DEL
-----
pg_catalog      pg_class      16      7162
546      319      0      1      0
pg_catalog      pg_index      4      660
145      149      0      0      0
pg_catalog      pg_namespace      4      100
49      49      0      0      0

```

pg_catalog	pg_database	4	16
42 42	0 0 0		
pg_catalog	pg_am	13	13
0 0	0 0 0		
sys	edb\$snap	1	9
0 0	1 0 0		
pg_catalog	pg_authid	1	1
25 25	0 0 0		
sys	edb\$session_wait_history	0	0
0 0	50 0 0		
sys	edb\$session_waits	0	0
0 0	2 0 0		
sys	edb\$stat_all_indexes	0	0
0 0	165 0 0		
DATA from pg_statio_all_tables			
SCHEMA	RELATION	HEAP	HEAP
IDX	TOAST	TIDX	TIDX
HIT	ICACHE	READ	HIT
HIT	HIT	HIT	HIT

pg_catalog	pg_class	0	539
1117 0	0 0	0	0
enterprisedb	pgbench_accounts	48	485
778 0	0 0	0	0
pg_catalog	pg_attribute	0	447
867 0	0 0	0	0
enterprisedb	pgbench_branches	0	439
114 0	0 0	0	0
enterprisedb	pgbench_tellers	0	357
112 0	0 0	0	0
pg_catalog	pg_statistic	1	293
441 0	0 0	0	0
pg_catalog	pg_index	0	159
171 0	0 0	0	0
pg_catalog	pg_opclass	0	145
68 0	0 0	0	0
pg_catalog	pg_proc	0	135
294 0	0 0	0	0
pg_catalog	pg_type	0	103
322 0	0 0	0	0
DATA from pg_stat_all_indexes			
SCHEMA	RELATION	INDEX	
IDX SCAN	IDX TUP READ	IDX TUP	FETCH

pg_catalog	pg_attribute		
pg_attribute_relid_attnum_index	427	907	907
pg_catalog	pg_class		pg_class_relname_nsp_index
289 62	62		
pg_catalog	pg_class		pg_class_oid_index
257 257	257		
pg_catalog	pg_statistic		
pg_statistic_relid_att_inh_index	207	196	196
enterprisedb	pgbench_accounts		pgbench_accounts_pkey
200 255	200		

```

pg_catalog      pg_cast      pg_cast_source_target_index
199      50      50
pg_catalog      pg_proc      pg_proc_oid_index
116      116      116
pg_catalog      edb_partition      edb_partition_partrelid_index
112      0      0
pg_catalog      edb_policy      edb_policy_object_name_index
112      0      0
enterprisedb    pgbench_branches      pgbench_branches_pkey
101      110      0

```

DATA from pg_statio_all_indexes

SCHEMA		RELATION		INDEX	
IDX	BLKS READ	IDX	BLKS HIT	IDX	BLKS ICACHE HIT

pg_catalog		pg_attribute			
pg_attribute_relid_attnum_index			0	867	0
enterprisedb		pgbench_accounts		pgbench_accounts_pkey	
1	778		0		
pg_catalog		pg_class		pg_class_relname_nsp_index	
0	590		0		
pg_catalog		pg_class		pg_class_oid_index	
0	527		0		
pg_catalog		pg_statistic			
pg_statistic_relid_att_inh_index			0	441	0
sys		edb\$stat_all_indexes		edb\$stat_idx_pk	
1	332		0		
sys		edb\$statio_all_indexes		edb\$statio_idx_pk	
1	332		0		
pg_catalog		pg_proc		pg_proc_oid_index	
0	244		0		
sys		edb\$stat_all_tables		edb\$stat_tab_pk	
0	241		0		
sys		edb\$statio_all_tables		edb\$statio_tab_pk	
0	241		0		

System Wait Information

WAIT NAME	COUNT	WAIT TIME	% WAIT

query plan	0	0.000407	100.00
db file read	0	0.000000	0.00

Database Parameters from postgresql.conf

PARAMETER		SETTING
CONTEXT	MINVAL	MAXVAL

allow_system_table_mods		off
postmaster		
application_name		psql
user		
archive_command		(disabled)
sighup		
archive_mode		off
postmaster		
archive_timeout		0
sighup	0	2147483647
array_nulls		on
user		

```

authentication_timeout      60
sighup      1      600
autovacuum      on
sighup
autovacuum_analyze_scale_factor  0.1
sighup      0      100
autovacuum_analyze_threshold    50
sighup      0      2147483647
autovacuum_freeze_max_age      200000000
postmaster 100000000  2000000000
autovacuum_max_workers      3
postmaster 1      8388607
autovacuum_naptime      60
sighup      1      2147483
autovacuum_vacuum_cost_delay    20
sighup      -1      100
autovacuum_vacuum_cost_limit   -1
sighup      -1      10000
autovacuum_vacuum_scale_factor  0.2
sighup      0      100
autovacuum_vacuum_threshold    50
sighup      0      2147483647
backslash_quote      safe_encoding
user
bgwriter_delay      200
sighup      10      10000
bgwriter_lru_maxpages      100
sighup      0      1000
bgwriter_lru_multiplier      2
sighup      0      10
block_size      8192
internal 8192      8192
bonjour      off
postmaster
bonjour_name
postmaster
bytea_output      hex
user
check_function_bodies      on
user
checkpoint_completion_target    0.5
sighup      0      1
checkpoint_segments      64
sighup      1      2147483647
checkpoint_timeout      300
sighup      30      3600
checkpoint_warning      30
sighup      0      2147483647
client_encoding      UTF8
user
client_min_messages      notice
user
commit_delay      0
user      0      100000
commit_siblings      5
user      0      1000
config_file      /opt/PostgresPlus/9.2AS/data/postgresql.
postmaster
constraint_exclusion      partition
user
cpu_index_tuple_cost      0.005
user      0      1.79769e+308

```



```

cpu_operator_cost          0.0025
user      0                1.79769e+308
cpu_tuple_cost             0.01
user      0                1.79769e+308
cursor_tuple_fraction      0.1
user      0                1
data_directory              /opt/PostgresPlus/9.2AS/data
postmaster
DateStyle                  Redwood, SHOW_TIME
user
db_dialect                  redwood
user
dbms_alert.max_alerts      100
postmaster 0                500
dbms_pipe.total_message_buffer 30
postmaster 30              262144
db_user_namespace          off
sighup
deadlock_timeout           1000
superuser 1                2147483647
debug_assertions            off
user
debug_pretty_print          on
user
debug_print_parse           off
user
debug_print_plan            off
user
debug_print_rewritten        off
user
default_heap_fillfactor    100
user      10                100
default_statistics_target   100
user      1                10000
default_tablespace          user
default_text_search_config  pg_catalog.english
user
default_transaction_deferrable off
user
default_transaction_isolation read committed
user
default_transaction_read_only off
user
default_with_oids           off
user
default_with_rowids          off
user
dynamic_library_path        $libdir
superuser
edb_audit
sighup
edb_audit_connect            failed
sighup
edb_audit_directory          edb_audit
sighup
edb_audit_disconnect         none
sighup
edb_audit_filename           audit-%Y%m%d_%H%M%S
sighup
edb_audit_rotation_day       every
sighup

```

```

edb_audit_rotation_seconds      0
sighup      0      2147483647
edb_audit_rotation_size      0
sighup      0      5000
edb_audit_statement      ddl, error
sighup
edb_connectby_order      on
user
edb_dynatune      66
postmaster  0      100
edb_dynatune_profile      oltp
postmaster
edb_enable_icache      off
postmaster
edb_icache_compression_level      6
superuser   0      9
edb_icache_servers
sighup
edb_redwood_date      on
user
edb_redwood_strings      on
user
edb_stmt_level_tx      off
user
effective_cache_size      34277
user      1      2147483647
effective_io_concurrency      1
user      0      1000
enable_bitmapscan      on
user
enable_hashagg      on
user
enable_hashjoin      on
user
enable_hints      on
user
enable_indexonlyscan      on
user
enable_indexscan      on
user
enable_material      on
user
enable_mergejoin      on
user
enable_nestloop      on
user
enable_seqscan      on
user
enable_sort      on
user
enable_tidscan      on
user
escape_string_warning      on
user
event_source      PostgreSQL
postmaster
exit_on_error      off
user
external_pid_file
postmaster
extra_float_digits      0
user      -15      3

```

```

from_collapse_limit      8
user      1      2147483647
fsync                      on
sighup
full_page_writes          on
sighup
geqo                      on
user
geqo_effort                5
user      1      10
geqo_generations          0
user      0      2147483647
geqo_pool_size            0
user      0      2147483647
geqo_seed                 0
user      0      1
geqo_selection_bias       2
user      1.5      2
geqo_threshold            12
user      2      2147483647
gin_fuzzy_search_limit    0
user      0      2147483647
hba_file                   /opt/PostgresPlus/9.2AS/data/pg_hba.conf
postmaster
hot_standby                off
postmaster
hot_standby_feedback      off
sighup
ident_file                 /opt/PostgresPlus/9.2AS/data/pg_ident.co
postmaster
ignore_system_indexes     off
backend
integer_datetimes         on
internal
IntervalStyle             postgres
user
join_collapse_limit      8
user      1      2147483647
krb_caseins_users        off
sighup
krb_server_keyfile        FILE:/home/edb/AS92/edb-postgres/inst/et
sighup
krb_srvname               postgres
sighup
lc_collate                en_US.UTF-8
internal
lc_ctype                  en_US.UTF-8
internal
lc_messages               en_US.UTF-8
superuser
lc_monetary               en_US.UTF-8
user
lc_numeric                en_US.UTF-8
user
lc_time                   en_US.UTF-8
user
listen_addresses          *
postmaster
local_preload_libraries
backend
lo_compat_privileges      off
superuser

```

```

log_autovacuum_min_duration      -1
sighup      -1      2147483647
log_checkpoints                  off
sighup
log_connections                  off
backend
log_destination                  stderr
sighup
log_directory                    pg_log
sighup
log_disconnections              off
backend
log_duration                    off
superuser
log_error_verbosity              default
superuser
log_executor_stats              off
superuser
log_file_mode                    0600
sighup      0      511
log_filename                    enterprisedb-%Y-%m-%d_%H%M%S.log
sighup
logging_collector                on
postmaster
log_hostname                    off
sighup
log_line_prefix                  %t
sighup
log_lock_waits                  off
superuser
log_min_duration_statement       -1
superuser      -1      2147483647
log_min_error_statement         error
superuser
log_min_messages                warning
superuser
log_parser_stats                off
superuser
log_planner_stats               off
superuser
log_rotation_age                1440
sighup      0      35791394
log_rotation_size               10240
sighup      0      2097151
log_statement                    none
superuser
log_statement_stats             off
superuser
log_temp_files                  -1
superuser      -1      2147483647
log_timezone                    US/Eastern
sighup
log_truncate_on_rotation        off
sighup
maintenance_work_mem            36871
user      1024      2097151
max_connections                 100
postmaster 1      8388607
max_files_per_process           1000
postmaster 25      2147483647
max_function_args               256
internal      256      256

```

```

max_identifier_length      63
internal      63          63
max_index_keys            32
internal      32          32
max_locks_per_transaction  64
postmaster  10          2147483647
max_pred_locks_per_transaction  64
postmaster  10          2147483647
max_prepared_transactions  0
postmaster  0          8388607
max_stack_depth          2048
superuser    100          2097151
max_standby_archive_delay  30000
sighup       -1          2147483647
max_standby_streaming_delay  30000
sighup       -1          2147483647
max_wal_senders           0
postmaster  0          8388607
odbc_lib_path
postmaster
optimizer_mode            choose
user
oracle_home
postmaster
password_encryption        on
user
plpgsql.variable_conflict  error
superuser
port                      5444
postmaster  1          65535
post_auth_delay           0
backend      0          2147483647
pre_auth_delay            0
sighup       0          60
qreplace_function
superuser
quote_all_identifiers      off
user
random_page_cost           4
user      0          1.79769e+308
replication_timeout        60000
sighup     0          2147483647
restart_after_crash        on
sighup
search_path                "$user",public
user
segment_size              131072
internal    131072      131072
seq_page_cost              1
user      0          1.79769e+308
server_encoding            UTF8
internal
server_version             9.2.0.0
internal
server_version_num         90200
internal    90200      90200
session_replication_role   origin
superuser
shared_buffers            23540
postmaster  16          1073741823
shared_preload_libraries   $libdir/dbms_pipe,$libdir/edb_gen,$libdi
postmaster

```

```

sql_inheritance          on
user
ssl                      off
postmaster
ssl_ca_file
postmaster
ssl_cert_file            server.crt
postmaster
ssl_ciphers              ALL:!ADH:!LOW:!EXP:!MD5:@STRENGTH
postmaster
ssl_crl_file
postmaster
ssl_key_file             server.key
postmaster
ssl_renegotiation_limit  524288
user      0              2097151
standard_conforming_strings on
user
statement_timeout        0
user      0              2147483647
stats_temp_directory      pg_stat_tmp
sighup
superuser_reserved_connections 3
postmaster 0              8388607
synchronize_seqscans      on
user
synchronous_commit        on
user
synchronous_standby_names
sighup
syslog_facility           local0
sighup
syslog_ident              postgres
sighup
tcp_keepalives_count      0
user      0              2147483647
tcp_keepalives_idle       0
user      0              2147483647
tcp_keepalives_interval   0
user      0              2147483647
temp_buffers              1024
user      100            1073741823
temp_file_limit            -1
superuser -1              2147483647
temp_tablespace
user
timed_statistics          on
user
TimeZone                  US/Eastern
user
timezone_abbreviations    Default
user
trace_hints               off
user
trace_notify              off
user
trace_recovery_messages    log
sighup
trace_sort                off
user
track_activities          on
superuser

```

```

track_activity_query_size      1024
postmaster 100                102400
track_counts                   on
superuser
track_functions                none
superuser
track_io_timing                off
superuser
transaction_deferrable         off
user
transaction_isolation           read committed
user
transaction_read_only           off
user
transform_null_equals           off
user
unix_socket_directory
postmaster
unix_socket_group
postmaster
unix_socket_permissions        0777
postmaster 0                  511
update_process_title            on
superuser
vacuum_cost_delay                0
user 0                        100
vacuum_cost_limit                200
user 1                        10000
vacuum_cost_page_dirty           20
user 0                        10000
vacuum_cost_page_hit             1
user 0                        10000
vacuum_cost_page_miss            10
user 0                        10000
vacuum_defer_cleanup_age         0
sighup 0                      1000000
vacuum_freeze_min_age            50000000
user 0                        1000000000
vacuum_freeze_table_age          150000000
user 0                        2000000000
wal_block_size                  8192
internal 8192                  8192
wal_buffers                      735
postmaster -1                  2147483647
wal_keep_segments                0
sighup 0                        2147483647
wal_level                        minimal
postmaster
wal_receiver_status_interval     10
sighup 0                        2147483
wal_segment_size                 2048
internal 2048                  2048
wal_sync_method                  fdatasync
sighup
wal_writer_delay                 200
sighup 1                        10000
work_mem                        3716
user 64                        2097151
xmlbinary                        base64
user
xmloption                        content
user

```

```
zero_damaged_pages      off
superuser
(412 rows)
```

4.3.2 stat_db_rpt()

The signature is:

```
stat_db_rpt(beginning_id, ending_id)
```

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

The following example demonstrates the `stat_db_rpt()` function:

```
SELECT * FROM stat_db_rpt(9, 10);
```

stat_db_rpt							
DATA from pg_stat_database							
DATABASE	NUMBACKENDS	XACT COMMIT	XACT ROLLBACK	BLKS READ	BLKS HIT		
	BLKS ICACHE HIT	HIT RATIO	ICACHE HIT RATIO				
edb	1	21	0	92928	101217		
	301	52.05	0.15				

4.3.3 stat_tables_rpt()

The signature is:

```
function_name(beginning_id, ending_id, top_n, scope)
```

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

`top_n`

`top_n` represents the number of rows to return

`scope`

`scope` determines which tables the function returns statistics about. Specify `SYS`, `USER` or `ALL`:

- `SYS` indicates that the function should return information about system defined tables. A table is considered a system table if it is stored in one of the following schemas: `pg_catalog`, `information_schema`, `sys`, or `dbo`.
- `USER` indicates that the function should return information about user-defined tables.
- `ALL` specifies that the function should return information about all tables.

The following code sample demonstrates the `stat_tables_rpt()` function:

```
SELECT * FROM stat_tables_rpt(18, 19, 10, 'ALL');
```

stat_tables_rpt														

DATA from pg_stat_all_tables ordered by seq scan														
SCHEMA	RELATION	SEQ	SCAN	REL	TUP	READ	IDX	SCAN	IDX	TUP	READ	INS	UPD	DEL

pg_catalog	pg_class	8					78		65			0	0	0
pg_catalog	pg_index	4					23		28			0	0	0
pg_catalog	pg_namespace	4					1		1			0	0	0
pg_catalog	pg_database	3					0		0			0	0	0
pg_catalog	pg_authid	2					0		0			0	0	0
sys	edb\$snap	1					0		0			1	0	0
public	accounts	0					0		0			0	0	0
public	branches	0					0		0			0	0	0
sys	edb\$session_wait_history	0					0		0			25	0	0
sys	edb\$session_waits	0					0		0			10	0	0

```
DATA from pg_stat_all_tables ordered by rel tup read
```

SCHEMA	RELATION	SEQ	SCAN	REL	TUP	READ	IDX	SCAN	IDX	TUP	READ	INS	UPD	DEL
pg_catalog	pg_class	8		2952			78		65			0	0	0
pg_catalog	pg_index	4		448			23		28			0	0	0
pg_catalog	pg_namespace	4		76			1		1			0	0	0
sys	edb\$snap	1		15			0		0			1	0	0
pg_catalog	pg_database	3		6			0		0			0	0	0
pg_catalog	pg_authid	2		1			0		0			0	0	0
public	accounts	0		0			0		0			0	0	0
public	branches	0		0			0		0			0	0	0
sys	edb\$session_wait_history	0		0			0		0			25	0	0
sys	edb\$session_waits	0		0			0		0			10	0	0
(29 rows)														

4.3.4 statio_tables_rpt()

The signature is:

```
statio_tables_rpt(beginning_id, ending_id, top_n, scope)
```

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

`top_n`

`top_n` represents the number of rows to return

`scope`

`scope` determines which tables the function returns statistics about. Specify `sys`, `USER` or `ALL`.

- **SYS** indicates that the function should return information about system defined tables. A table is considered a system table if it is stored in one of the following schemas: `pg_catalog`, `information_schema`, `sys`, or `dbo`.
- **USER** indicates that the function should return information about user-defined tables.
- **ALL** specifies that the function should return information about all tables.

The following example demonstrates the `statio_tables_rpt()` function:

```
edb=# SELECT * FROM statio_tables_rpt(9, 10, 10, 'SYS');
```

statio_tables_rpt							

DATA from pg_statio_all_tables							
SCHEMA	RELATION	HEAP READ	HEAP HIT	HEAP ICACHE HIT	IDX READ	IDX HIT	
	IDX ICACHE HIT	TOAST READ	TOAST HIT	TOAST ICACHE HIT	TIDX READ	TIDX HIT	TIDX ICACHE HIT

public	pgbench_accounts	92766	67215	288	59	32126	
pg_catalog	pg_class	0	0	296	0	0	16
sys	edb\$stat_all_indexes	8	125	0	4	233	
sys	edb\$statio_all_index	8	125	0	4	233	
sys	edb\$stat_all_tables	6	91	0	2	174	
sys	edb\$statio_all_table	6	91	0	2	174	
pg_catalog	pg_namespace	3	72	0	0	0	0
sys	edb\$session_wait_his	1	24	0	4	47	
pg_catalog	pg_opclass	3	13	0	2	0	
pg_catalog	pg_trigger	0	12	0	1	15	
(16 rows)							

4.3.5 stat_indexes_rpt()

The signature is:

```
stat_indexes_rpt(beginning_id, ending_id, top_n, scope)
```

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

`top_n`

`top_n` represents the number of rows to return

`scope`

`scope` determines which tables the function returns statistics about. Specify `SYS`, `USER` or `ALL`:

- `SYS` indicates that the function should return information about system defined tables. A table is considered a system table if it is stored in one of the following schemas: `pg_catalog`, `information_schema`, `sys`, or `dbo`.
- `USER` indicates that the function should return information about user-defined tables.
- `ALL` specifies that the function should return information about all tables.

The following code sample demonstrates the `stat_indexes_rpt()` function:

```
edb=# SELECT * FROM stat_indexes_rpt(9, 10, 10, 'ALL');
```

stat_indexes_rpt				

DATA from pg_stat_all_indexes				
SCHEMA	RELATION	INDEX		
		IDX SCAN	IDX TUP READ	IDX TUP FETCH

pg_catalog	pg_cast	pg_cast_source_target_index		
		30	7	7
pg_catalog	pg_class	pg_class_oid_index		
		15	15	15
pg_catalog	pg_trigger	pg_trigger_tgrelid_tgname_index		
		12	12	12
pg_catalog	pg_attribute	pg_attribute_relid_attnum_index		
		7	31	31
pg_catalog	pg_statistic	pg_statistic_relid_att_index		
		7	0	0

```

pg_catalog    pg_database    pg_database_oid_index
              5              5              5
pg_catalog    pg_proc      pg_proc_oid_index
              5              5              5
pg_catalog    pg_operator   pg_operator_oprname_l_r_n_index
              3              1              1
pg_catalog    pg_type       pg_type_typname_nsp_index
              3              1              1
pg_catalog    pg_amop       pg_amop_opr_fam_index
              2              3              3
(14 rows)

```

4.3.6 statio_indexes_rpt()

The signature is:

```
statio_indexes_rpt(beginning_id, ending_id, top_n, scope)
```

Parameters

`beginning_id`

`beginning_id` is an integer value that represents the beginning session identifier.

`ending_id`

`ending_id` is an integer value that represents the ending session identifier.

`top_n`

`top_n` represents the number of rows to return

`scope`

`scope` determines which tables the function returns statistics about. Specify `SYS`, `USER` or `ALL`:

- `SYS` indicates that the function should return information about system defined tables. A table is considered a system table if it is stored in one of the following schemas: `pg_catalog`, `information_schema`, `sys`, or `dbo`.
- `USER` indicates that the function should return information about user-defined tables.
- `ALL` specifies that the function should return information about all tables.

The following example demonstrates the `statio_indexes_rpt()` function:

```
edb=# SELECT * FROM statio_indexes_rpt(9, 10, 10, 'SYS');
```

statio_indexes_rpt					
DATA from pg_statio_all_indexes					
SCHEMA	RELATION	INDEX	IDX BLKS READ	IDX BLKS HIT	IDX BLKS ICACHE HIT
public	pgbench_accounts	pgbench_accounts_pkey	59	32126	9
sys	edb\$stat_all_indexes	edb\$stat_idx_pk	4	233	0
sys	edb\$statio_all_indexes	edb\$statio_idx_pk	4	233	0
sys	edb\$stat_all_tables	edb\$stat_tab_pk	2	174	0
sys	edb\$statio_all_tables	edb\$statio_tab_pk	2	174	0
sys	edb\$session_wait_history	session_waits_hist_pk	4	47	0
pg_catalog	pg_cast	pg_cast_source_target_index	1	29	0
pg_catalog	pg_trigger	pg_trig_tgrelid_tgname_index	1	15	0
pg_catalog	pg_class	pg_class_oid_index	1	14	0
pg_catalog	pg_statistic	pg_statistic_relid_att_index	2	12	0

(14 rows)

4.4 Performance Tuning Recommendations

To use DRITA reports for performance tuning, review the top five events in a given report, looking for any event that takes a disproportionately large percentage of resources. In a streamlined system, user I/O will probably make up the largest number of waits. Waits should be evaluated in the context of CPU usage and total time; an event may not be significant if it takes 2 minutes out of a total measurement interval of 2 hours, if the rest of the time is consumed by CPU time. The component of response time (CPU "work" time or other "wait" time) that consumes the highest percentage of overall time should be evaluated.

When evaluating events, watch for:

Event type	Description
Checkpoint waits	Checkpoint waits may indicate that checkpoint parameters need to be adjusted, (<code>checkpoint_segments</code> and <code>checkpoint_timeout</code>).
WAL-related waits	WAL-related waits may indicate <code>wal_buffers</code> are under-sized.
SQL Parse waits	If the number of waits is high, try to use prepared statements.
db file random reads	If high, check that appropriate indexes and statistics exist.
db file random writes	If high, may need to decrease <code>bgwriter_delay</code> .
btree random lock acquires	May indicate indexes are being rebuilt. Schedule index builds during less active time.

Performance reviews should also include careful scrutiny of the hardware, the operating system, the network and the application SQL statements.

4.5 Event Descriptions

Event Name	Description
add in shmem lock acquire	Obsolete/unused
bgwriter communication lock acquire	The bgwriter (background writer) process has waited for the short-term lock that synchronizes messages between the bgwriter and a backend process.
btree vacuum lock acquire	The server has waited for the short-term lock that synchronizes access to the next available vacuum cycle ID.
buffer free list lock acquire	The server has waited for the short-term lock that synchronizes access to the list of free buffers (in shared memory).
checkpoint lock acquire:	A server process has waited for the short-term lock that prevents simultaneous checkpoints.
checkpoint start lock acquire	The server has waited for the short-term lock that synchronizes access to the bgwriter checkpoint schedule.
clog control lock acquire	The server has waited for the short-term lock that synchronizes access to the commit log.
control file lock acquire	The server has waited for the short-term lock that synchronizes write access to the control file (this should usually be a low number).
db file extend	A server process has waited for the operating system while adding a new page to the end of a file.
db file read	A server process has waited for the completion of a read (from disk).
db file write	A server process has waited for the completion of a write (to disk).
db file sync	A server process has waited for the operating system to flush all changes to disk.
first buf mapping lock acquire	The server has waited for a short-term lock that synchronizes access to the shared-buffer mapping table.
freespace lock acquire	The server has waited for the short-term lock that synchronizes access to the freespace map.
Infinite Cache read	The server has waited for an Infinite Cache read request.
Infinite Cache write	The server has waited for an Infinite Cache write request.
lwlock acquire	The server has waited for a short-term lock that has not been described elsewhere in this section.
multi xact gen lock acquire	The server has waited for the short-term lock that synchronizes access to the next available multi-transaction ID (when a SELECT...FOR SHARE statement executes).
multi xact member lock acquire	The server has waited for the short-term lock that synchronizes access to the multi-transaction member file (when a SELECT...FOR SHARE statement executes).
multi xact offset lock acquire	The server has waited for the short-term lock that synchronizes access to the multi-transaction offset file (when a SELECT...FOR SHARE statement executes).
oid gen lock acquire	The server has waited for the short-term lock that synchronizes access to the next available OID (object ID).
query plan	The server has computed the execution plan for a SQL statement.
rel cache init lock acquire	The server has waited for the short-term lock that prevents simultaneous relation-cache loads/unloads.
shmem index lock acquire	The server has waited for the short-term lock that synchronizes access to the shared-memory map.
sinval lock acquire	The server has waited for the short-term lock that synchronizes access to the cache invalidation state.
sql parse	The server has parsed a SQL statement.

subtrans control lock acquire	The server has waited for the short-term lock that synchronizes access to the subtransaction log.
tablespace create lock acquire	The server has waited for the short-term lock that prevents simultaneous CREATE TABLESPACE or DROP TABLESPACE commands.
two phase state lock acquire	The server has waited for the short-term lock that synchronizes access to the list of prepared transactions.
wal insert lock acquire	The server has waited for the short-term lock that synchronizes write access to the write-ahead log. A high number may indicate that WAL buffers are sized too small.
wal write lock acquire	The server has waited for the short-term lock that synchronizes write-ahead log flushes.
wal file sync	The server has waited for the write-ahead log to sync to disk (related to the wal_sync_method parameter which, by default, is 'fsync' - better performance can be gained by changing this parameter to open_sync).
wal flush	The server has waited for the write-ahead log to flush to disk.
wal write	The server has waited for a write to the write-ahead log buffer (expect this value to be high).
xid gen lock acquire	The server has waited for the short-term lock that synchronizes access to the next available transaction ID.

4.6 Catalog Views

The following DRITA catalog views provide access to performance information relating to system waits.

4.6.1 edb\$system_waits

The `edb$system_waits` view summarizes the number of waits and the total wait time per session for each wait named. It also displays the average and max wait times.

`edb$system_waits` summarizes the following information:

Column	Type	Modifiers	Definition
<code>edb_id</code>	numeric		identifier
<code>dbname</code>	text		database name
<code>wait_name</code>	text		name of the event
<code>wait_count</code>	numeric		number of times the event occurs
<code>avg_wait</code>	numeric(50,6)		average wait time in microseconds
<code>max_wait</code>	numeric		maximum wait time in microseconds
<code>total_wait</code>	numeric		total wait time in microseconds
<code>wait_name</code>	text		name of the event

The following example shows the result of a `SELECT` statement on the `edb$system_waits` view:

```
select * from sys.edb$system_waits;
```

edb_id	dbname	wait_name	wait_count	avg_wait	max_wait	totalwait
1	edb	db fileread	301	0.011516	0.629986	2.742500
1	edb	wal flush	26	0.010364	0.085380	0.269452
1	edb	wal write	26	0.010355	0.085371	0.269232
1	edb	query plan	277	0.001367	0.049425	0.192442
2	edb	wal flush	28	0.040443	0.095150	0.431984
2	edb	wal write	28	0.040434	0.095093	0.431698
2	edb	query plan	299	0.001479	0.049425	0.262596

4.6.2 edb\$session_waits

The `edb$session_waits` view summarizes the number of waits and the total wait time per session for each wait named and identified by backend ID. It also displays the average and max wait times. `edb$session_waits` summarizes the following information:

Column	Type	Modifiers	Definition
<code>backend_id</code>	bigint		session identifier
<code>wait_count</code>	bigint		number of times the event occurs
<code>avg_wait_time</code>	numeric		average wait time in microseconds
<code>max_wait_time</code>	numeric(50,6)		maximum wait time in microseconds

total_wait_time	numeric(50,6)		total wait time in microseconds
wait_name	text		name of the event

The following code sample shows the result of a `SELECT` statement on the `edb$session_waits` view:

```
SELECT * FROM sys.edb$session_waits;
```

edb_id	dbname	backend_id	wait_name	wait_count	avg_wait_time
max_wait_time	total_wait_time	username	current_query		
1	edb	22935	db file read	175	0.008399
0.629986		1.469887	enterprisedb	<IDLE>	
1	edb	22988	db file read	116	0.009556
0.040627		1.108438	enterprisedb	select * from edbsnap();	
1	edb	22988	wal flush	26	0.010364
0.085380		0.269452	enterprisedb	select * from edbsnap();	

(3 rows)

4.6.3 edb\$session_wait_history

The `edb$session_wait_history` view contains the last 25 wait events for each backend ID active during the session. The `edb$session_wait_history` view includes the following information:

Column	Type	Modifiers	Definition
edb_id	numeric		identifier
dbname	text		database name
backend_id	bigint		session identifier
seq	bigint		number between 1 and 25
wait_name	text		name of the event
elapsed	bigint		elapsed time in microseconds
p1	bigint		variable #1- meaning dependent on event
p2	bigint		variable #2- meaning dependent on event
p3	bigint		variable #3- meaning dependent on event

The following code sample shows the result of a `SELECT` statement on the `edb$session_wait_history` view:

```
SELECT * FROM sys.edb$session_wait_history;
```

edb_id	dbname	backend_id	seq	wait_name	elapsed	p1	p2	p3
1	edb	22935	1	query plan	54	0	0	0
1	edb	22935	2	db file read	1116	2689	8	1
1	edb	22935	3	db file read	983	1255	32	1
1	edb	22935	4	db file read	13717	2691	19	1
1	edb	22935	5	query plan	75	0	0	0
1	edb	22935	6	db file read	11053	1255	7	1
1	edb	22935	7	db file read	404	2689	4	1

(7 rows)

5 DBMS_PROFILER

The DBMS_PROFILER package collects and stores performance information about the PL/pgSQL and SPL statements that are executed during a profiling session; you can review the performance information in the tables and views provided by the profiler.

DBMS_PROFILER works by recording a set of performance-related counters and timers for each line of PL/pgSQL or SPL statement that executes within a profiling session. The counters and timers are stored in a table named SYS.PLSQL_PROFILER_DATA. When you complete a profiling session, DBMS_PROFILER will write a row to the performance statistics table for each line of PL/pgSQL or SPL code that executed within the session. For example, if you execute the following function:

```

1 - CREATE OR REPLACE FUNCTION getBalance(acctNumber INTEGER)
2 - RETURN NUMBER AS
3 -     result NUMBER;
4 - BEGIN
5 -     SELECT balance INTO result FROM acct WHERE id = acctNumber;
6 -
7 -     IF (balance IS NULL) THEN
8 -         DBMS_OUTPUT.PUT_LINE('Balance is null');
9 -     END IF;
10-
11-     RETURN result;
12- END;
```

DBMS_PROFILER adds one PLSQL_PROFILER_DATA entry for each line of code within the getBalance() function (including blank lines and comments). The entry corresponding to line 4 will show that the SELECT statement executed exactly one time; and required a very small amount of time to execute. On the other hand, the entry corresponding to line 8 will show that the call to DBMS_OUTPUT.PUT_LINE executed once or not at all (depending on the value for the balance column).

Some of the lines in this function contain no executable code (for example, 6, 9, and 10) so the performance statistics for those lines will always contain *zero* values.

To start a profiling session, invoke the DBMS_PROFILER.START_PROFILER function (or procedure). Once you've invoked START_PROFILER, Advanced Server will profile every PL/pgSQL or SPL function, procedure, trigger, or anonymous block that your session executes until you either stop or pause the profiler (by calling STOP_PROFILER or PAUSE_PROFILER).

It is important to note that when you start (or resume) the profiler, the profiler will only gather performance statistics for functions/procedures/triggers that *start* after the call to START_PROFILER (or RESUME_PROFILER).

While the profiler is active, Advanced Server records a large set of timers and counters in memory; when you invoke the `STOP_PROFILER` (or `FLUSH_DATA`) function/procedure, `DBMS_PROFILER` writes those timers and counters to a set of three tables:

- `SYS.PLSQL_PROFILER_RAWDATA`
Contains the performance counters and timers for each statement executed within the session.
- `SYS.PLSQL_PROFILER_RUNS`
Contains a summary of each run (aggregating the information found in `PLSQL_PROFILER_RAWDATA`).
- `SYS.PLSQL_PROFILER_UNITS`
Contains a summary of each code unit (function, procedure, trigger, or anonymous block) executed within a session.

In addition, `DBMS_PROFILER` defines a view, `SYS.PLSQL_PROFILER_DATA`, which contains a subset of the `PLSQL_PROFILER_RAWDATA` table (in a form that is compatible with Oracle's `DBMS_PROFILER` package).

Please note that a non-superuser may *gather* profiling information, but may not view that profiling information unless a superuser grants specific privileges on the profiling tables (stored in the `SYS` schema). This permits a non-privileged user to gather performance statistics without exposing information that the administrator may want to keep secret (i.e., PL/SQL code).

5.1 Querying the DBMS_PROFILER Tables and View

The following step-by-step example uses DBMS_PROFILER to retrieve performance information for procedures, functions, and triggers included in the sample data distributed with Advanced Server.

1. Open the EDB-PSQL command line, and establish a connection to the Advanced Server database. Use an EXEC statement to start the profiling session:

```
acctg=# EXEC dbms_profiler.start_profiler('profile list_emp');
EDB-SPL Procedure successfully completed
```

(Note: the call to `start_profiler()` includes a comment that DBMS_PROFILER associates with the profiler session).

2. Then, use an EXEC statement a call to the `list_emp` procedure:

```
acctg=# EXEC list_emp;
EMPNO      ENAME
-----
7369       SMITH
7499       ALLEN
7521       WARD
7566       JONES
7654       MARTIN
7698       BLAKE
7782       CLARK
7788       SCOTT
7839       KING
7844       TURNER
7876       ADAMS
7900       JAMES
7902       FORD
7934       MILLER
```

3. Stop the profiling session with a call to `dbms_profiler.stop_profiler`:

```
EDB-SPL Procedure successfully completed
acctg=# EXEC dbms_profiler.stop_profiler;
```

4. Start a new session with the `dbms_profiler.start_profiler` function (followed by a new comment):

```
EDB-SPL Procedure successfully completed
acctg=# EXEC dbms_profiler.start_profiler('profile get_dept_name
and emp_sal_trig');
EDB-SPL Procedure successfully completed
```

5. Invoke the `get_dept_name` function:

```
acctg=# SELECT emp_admin.get_dept_name(10);
       get_dept_name
-----
ACCOUNTING
(1 row)
```

6. Execute an UPDATE statement that causes a trigger to execute:

```
acctg=# UPDATE memp SET sal = 500 WHERE empno = 7902;
Updating employee 7902
..Old salary: 3000.00
..New salary: 500.00
..Raise      : -2500.00
User korrry updated employee(s) on 2011-03-11
UPDATE 1
```

7. Terminate the profiling session and flush the performance information to the profiling tables:

```
acctg=# EXEC dbms_profiler.stop_profiler;
```

8. Now, query the `plsql_profiler_runs` table to view a list of the profiling sessions, arranged by runid:

```
EDB-SPL Procedure successfully completed
acctg=# SELECT * FROM plsql_profiler_runs;
```

runid	related_run	run_owner	run_date	run_comment	run_total_time	run_system_info	run_comment1	spare1
1		korrry	11-MAR-11 12:19:04.439087	profile list_emp	4211			
2		korrry	11-MAR-11 12:19:53.655886	profile get_dept_name and emp_sal_trig	16950			

```
(2 rows)
```

9. Query the `plsql_profiler_units` table to view the amount of time consumed by each unit (each function, procedure, or trigger):

```
acctg=# SELECT * FROM plsql_profiler_units;
```

runid	unit_number	unit_type	unit_owner	unit_name	unit_timestamp	total_time	spare1	spare2
1	16895	PROCEDURE	korrry	list_emp()		4		
2	16911	FUNCTION	korrry	get_dept_name(p_deptno numeric)		0		

```

2 | 16908 | | korry | emp_sal_trig_memp()
| 2 | 16906 | | korry | user_audit_trig_memp()
| 15 |
(4 rows)

```

10. Query the plsql_profiler_rawdata table to view a list of the wait event counters and wait event times:

```

acctg=# SELECT * FROM plsql_profiler_rawdata;
runid | sourcecode
      | func_oid | line_number | exec_count | tuples_returned |
time_+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
--
1 |          | 16895 | 0 | 0 | 0 |
1 | v_empno   | 16895 | 1 | 0 | 0 |
   |          |      |   |   |   |
1 | v_ename   | 16895 | 2 | 0 | 0 |
   |          |      |   |   |   |
1 | CURSOR emp_cur IS
   | 16895 | 3 | 0 | 0 |
1 | SELECT empno, ename FROM memp ORDER BY empno;
   | 16895 | 4 | 0 | 0 |
1 | BEGIN
   | 16895 | 5 | 1 | 0 |
0.0 1 | OPEN emp_cur;
   | 16895 | 6 | 1 | 0 |
0.0 1 | DBMS_OUTPUT.PUT_LINE('EMPNO   ENAME');
   | 16895 | 7 | 1 | 0 |
0.0 1 | DBMS_OUTPUT.PUT_LINE('-----');
   | 16895 | 8 | 1 | 0 |
0.0 1 | LOOP
   | 16895 | 9 | 15 | 0 |
0.0 1 | FETCH emp_cur INTO v_empno, v_ename;
   | 16895 | 10 | 15 | 0 |
0.0 1 | EXIT WHEN emp_cur%NOTFOUND;
   | 16895 | 11 | 14 | 0 |
0.0 1 | DBMS_OUTPUT.PUT_LINE(v_empno || ' ' || v_ename);
   | 16895 | 12 | 0 | 0 |
1 | END LOOP;
   | 16895 | 13 | 1 | 0 |
2.

```


1		CLOSE emp_cur;				
		16895		14		0 0
1		END				
		16895		15		0 0
2		DECLARE				
		16908		0		0 0
2		sal_diff NUMBER;				
		16908		1		0 0
2		BEGIN				
		16908		2		0 0
2		IF INSERTING THEN				
		16908		3		1 0
5.						
2		DBMS_OUTPUT.PUT_LINE('Inserting employee ' :NEW.empno);				
		16908		4		0 0
2		DBMS_OUTPUT.PUT_LINE('..New salary: ' :NEW.sal);				
		16908		5		0 0
2		END IF;				
		16908		6		0 0
2		IF UPDATING THEN				
		16908		7		1 0
0.0						
2		sal_diff := :NEW.sal - :OLD.sal;				
		16908		8		1 0
0.0						
2		DBMS_OUTPUT.PUT_LINE('Updating employee ' :OLD.empno);				
		16908		9		1 0
0.0						
2		DBMS_OUTPUT.PUT_LINE('..Old salary: ' :OLD.sal);				
		16908		10		1 0
0.0						
2		DBMS_OUTPUT.PUT_LINE('..New salary: ' :NEW.sal);				
		16908		11		1 0
0.0						
2		DBMS_OUTPUT.PUT_LINE('..Raise : ' sal_diff);				
		16908		12		1 0
0.						
2		END IF;				
		16908		13		0 0
2		IF DELETING THEN				
		16908		14		1 0
3.						
2		DBMS_OUTPUT.PUT_LINE('Deleting employee ' :OLD.empno);				
		16908		15		0 0
2		DBMS_OUTPUT.PUT_LINE('..Old salary: ' :OLD.sal);				
		16908		16		0 0
2		END IF;				
		16908		17		0 0
2		END				
		16908		18		0 0

```

2 |          | 16911 |          0 |          0 |          0 |
2 |          v_dname          VARCHAR2(14);
  |          | 16911 |          1 |          0 |          0 |
2 | BEGIN
  |          | 16911 |          2 |          1 |          0 |
0.0
2 |          SELECT dname INTO v_dname FROM mdept WHERE deptno = p_deptno;
  |          | 16911 |          3 |          1 |          0 |
3.
2 |          RETURN v_dname;
  |          | 16911 |          4 |          0 |          0 |
2 | EXCEPTION
  |          | 16911 |          5 |          0 |          0 |
2 |          WHEN NO_DATA_FOUND THEN
  |          | 16911 |          6 |          0 |          0 |
2 |          DBMS_OUTPUT.PUT_LINE('Invalid department number ' ||
p_deptno);          | 16911 |          7 |          0 |          0 |
|
2 |          RETURN '';
  |          | 16911 |          8 |          0 |          0 |
2 | END
  |          | 16911 |          9 |          0 |          0 |
2 | DECLARE
  |          | 16906 |          0 |          0 |          0 |
2 |          v_action          VARCHAR2(24);
  |          | 16906 |          1 |          0 |          0 |
2 | BEGIN
  |          | 16906 |          2 |          0 |          0 |
2 |          IF INSERTING THEN
  |          | 16906 |          3 |          1 |          0 |
0.0
2 |          v_action := ' added employee(s) on ';
  |          | 16906 |          4 |          0 |          0 |
2 |          ELSIF UPDATING THEN
  |          | 16906 |          5 |          1 |          0 |
6.
2 |          v_action := ' updated employee(s) on ';
  |          | 16906 |          6 |          1 |          0 |
3.
2 |          ELSIF DELETING THEN
  |          | 16906 |          7 |          0 |          0 |
2 |          v_action := ' deleted employee(s) on ';
  |          | 16906 |          8 |          0 |          0 |
2 |          END IF;
  |          | 16906 |          9 |          0 |          0 |
2 |          DBMS_OUTPUT.PUT_LINE('User ' || USER || v_action ||
TO_CHAR(SYSDATE, 'YYYY-MM-DD')); | 16906 |          10 |          1 |
0 | 0.0

```

```

2 | END
      |      16906 |      11 |      0 |      0 |
(57 rows)

```

11. Query the `plsql_profiler_data` view to review an Oracle-compatible subset of the information found in `plsql_profiler_rawdata` table:

```

acctg=# SELECT * FROM plsql_profiler_data;
 runid | unit_number | line# | total_occur | total_time | min_time | max_time |
 spare1 | spare2 | spare3 | spare4
-----+-----+-----+-----+-----+-----+-----+
 1 |      16895 |      0 |           0 |           0 |           0 |           0 |
 1 |      16895 |      1 |           0 |           0 |           0 |           0 |
 1 |      16895 |      2 |           0 |           0 |           0 |           0 |
 1 |      16895 |      3 |           0 |           0 |           0 |           0 |
 1 |      16895 |      4 |           0 |           0 |           0 |           0 |
 1 |      16895 |      5 |           1 | 0.000322 | 0.000322 | 0.000322 |
 1 |      16895 |      6 |           1 | 0.000675 | 0.000675 | 0.000675 |
 1 |      16895 |      7 |           1 | 0.000115 | 0.000115 | 0.000115 |
 1 |      16895 |      8 |           1 | 0.001605 | 0.001605 | 0.001605 |
 1 |      16895 |      9 |          15 | 0.000183 | 5e-06 | 5e-05 |
 1 |      16895 |     10 |          15 | 0.000458 | 5e-06 | 0.000378 |
 1 |      16895 |     11 |          14 | 0.000831 | 2.9e-05 | 0.000448 |
 1 |      16895 |     12 |           0 |           0 |           0 |           0 |
 1 |      16895 |     13 |           1 | 2.2e-05 | 2.2e-05 | 2.2e-05 |
 1 |      16895 |     14 |           0 |           0 |           0 |           0 |
 1 |      16895 |     15 |           0 |           0 |           0 |           0 |
 2 |      16908 |      0 |           0 |           0 |           0 |           0 |
 2 |      16908 |      1 |           0 |           0 |           0 |           0 |
 2 |      16908 |      2 |           0 |           0 |           0 |           0 |
 2 |      16908 |      3 |           1 | 5.1e-05 | 5.1e-05 | 5.1e-05 |
 2 |      16908 |      4 |           0 |           0 |           0 |           0 |
 2 |      16908 |      5 |           0 |           0 |           0 |           0 |
 2 |      16908 |      6 |           0 |           0 |           0 |           0 |
 2 |      16908 |      7 |           1 | 0.000846 | 0.000846 | 0.000846 |

```

Postgres Plus Advanced Server Performance and Scalability Guide

2	16908	8	1	0.000129	0.000129	0.000129
2	16908	9	1	0.000239	0.000239	0.000239
2	16908	10	1	0.000163	0.000163	0.000163
2	16908	11	1	0.000158	0.000158	0.000158
2	16908	12	1	0.00011	0.00011	0.00011
2	16908	13	0	0	0	0
2	16908	14	1	3.4e-05	3.4e-05	3.4e-05
2	16908	15	0	0	0	0
2	16908	16	0	0	0	0
2	16908	17	0	0	0	0
2	16908	18	0	0	0	0
2	16911	0	0	0	0	0
2	16911	1	0	0	0	0
2	16911	2	1	0.000328	0.000328	0.000328
2	16911	3	1	3.8e-05	3.8e-05	3.8e-05
2	16911	4	0	0	0	0
2	16911	5	0	0	0	0
2	16911	6	0	0	0	0
2	16911	7	0	0	0	0
2	16911	8	0	0	0	0
2	16911	9	0	0	0	0
2	16906	0	0	0	0	0
2	16906	1	0	0	0	0
2	16906	2	0	0	0	0
2	16906	3	1	0.000138	0.000138	0.000138
2	16906	4	0	0	0	0
2	16906	5	1	6.6e-05	6.6e-05	6.6e-05
2	16906	6	1	3.2e-05	3.2e-05	3.2e-05
2	16906	7	0	0	0	0
2	16906	8	0	0	0	0
2	16906	9	0	0	0	0

2		16906		10		1		0.014618		0.014618		0.014618	
2		16906		11		0		0		0		0	

(57 rows)

5.2 DBMS_PROFILER Functions and Procedures

The DBMS_PROFILER package collects and stores performance information about the PL/pgSQL and SPL statements that are executed during a profiling session; use the functions and procedures listed below to control the profiling tool.

Table 5-1 DBMS_PROFILER Functions/Procedures

Function/Procedure	Function or Procedure	Return Type	Description
FLUSH_DATA	Function and Procedure	Status Code or Exception	Flushes performance data collected in the current session without terminating the session (profiling continues).
GET_VERSION (major OUT, minor OUT)	Procedure	n/a	Returns the version number of this package.
INTERNAL_VERSION_CHECK	Function	Status Code	Confirms that the current version of the profiler will work with the current database.
PAUSE_PROFILER	Function and Procedure	Status Code or Exception	Pause data collection.
RESUME_PROFILER	Function and Procedure	Status Code or Exception	Resume data collection.
START_PROFILER[run_comment, run_comment1, run_number OUT]	Functions and Procedures	Status Code or Exception	Start data collection.
STOP_PROFILER	Function and Procedure	Status Code or Exception	Stop data collection and flush performance data to PLSQL_PROFILER_RAWDATA.

Return Values

The functions within the DBMS_PROFILER package return a status code to indicate success or failure; the DBMS_PROFILER procedures raise an exception only if they encounter a failure. The status codes and messages returned by the functions, and the exceptions raised by the procedures are listed in the table below.

Status Code	Message	Exception	Description
-1	error version	version_mismatch	The profiler version and the database are incompatible.
0	success	n/a	The operation completed successfully.
1	error_param	profiler_error	The operation received an incorrect parameter.
2	error_io	profiler_error	The data flush operation has failed.

5.2.1 FLUSH_DATA

The FLUSH_DATA procedure or function flushes the data collected in the current session without terminating the profiler session. The data is flushed to the tables listed in Section 6.3 of the Postgres Plus Advanced Server Performance Features Guide. The signature of the FLUSH_DATA function is:

```
DBMS_PROFILER.FLUSH_DATA
RETURN INTEGER;
```

The signature of the `FLUSH_DATA` procedure is:

```
DBMS_PROFILER.FLUSH_DATA;
```

5.2.2 GET_VERSION

The `GET_VERSION` procedure returns the version of `DBMS_PROFILER`. The procedure signature is:

```
DBMS_PROFILER.GET_VERSION( major OUT INTEGER
                           minor OUT INTEGER);
```

Parameters

`major`

The major version number of `DBMS_PROFILER`.

`minor`

The minor version number of `DBMS_PROFILER`.

5.2.3 INTERNAL_VERSION_CHECK

The `INTERNAL_VERSION_CHECK` function confirms that the current version of `DBMS_PROFILER` will work with the current database. The function signature is:

```
DBMS_PROFILER.INTERNAL_VERSION_CHECK
RETURN INTEGER;
```

5.2.4 PAUSE_PROFILER

The `PAUSE_PROFILER` function or procedure pauses a profiling session. The function signature is:

```
DBMS_PROFILER.PAUSE_PROFILER
RETURN INTEGER;
```

The signature of the `PAUSE_PROFILER` procedure is:

```
DBMS_PROFILER.PAUSE_PROFILER;
```

5.2.5 RESUME_PROFILER

The `RESUME_PROFILER` function or procedure resumes a paused profiling session. The function signature is:

```
DBMS_PROFILER.RESUME_PROFILER
    RETURN INTEGER;
```

The signature of the `RESUME_PROFILER` procedure is:

```
DBMS_PROFILER.RESUME_PROFILER;
```

5.2.6 START_PROFILER

The `START_PROFILER` function or procedure starts a data collection session. The `START_PROFILER` function has two forms:

```
DBMS_PROFILER.START_PROFILER (
    run_comment IN TEXT := sysdate,
    run_comment1 IN TEXT := '',
    run_number OUT INTEGER)
    RETURN INTEGER;
```

```
DBMS_PROFILER.START_PROFILER (
    run_comment IN TEXT := sysdate,
    run_comment1 IN TEXT := '')
    RETURN INTEGER;
```

The `START_PROFILER` procedure has two forms:

```
DBMS_PROFILER.START_PROFILER (
    run_comment IN TEXT := sysdate,
    run_comment1 IN TEXT := '');
```

```
DBMS_PROFILER.START_PROFILER (
    run_comment IN TEXT := sysdate,
    run_comment1 IN TEXT := '',
    run_number OUT INTEGER);
```

Parameters

`run_comment`

A user-defined comment for the profiler session; the default value is `sysdate`.

`run_comment1`

An additional user-defined comment for the profiler session; the default value is ''.

`run_number`

The session number of the profiler session.

5.2.7 STOP_PROFILER

The `STOP_PROFILER` function or procedure stops a profiling session and flushes the performance information to the `DBMS_PROFILER` tables and view. The `STOP_PROFILER` function signature is:

```
DBMS_PROFILER.STOP_PROFILER
RETURN INTEGER;
```

The signature of the `START_PROFILER` procedure is:

```
DBMS_PROFILER.STOP_PROFILER;
```

5.3 DBMS_PROFILER - Reference

The Advanced Server installer creates the following tables and views that you can query to review PL/SQL performance profile information:

Table Name	Description
PLSQL_PROFILER_RUNS	Table containing information about all profiler runs, organized by runid.
PLSQL_PROFILER_UNITS	Table containing information about all profiler runs, organized by unit.
PLSQL_PROFILER_DATA	Oracle-compatible view containing performance statistics.
PLSQL_PROFILER_RAWDATA	Table containing the Oracle-compatible performance statistics and the extended performance statistics for DRITA counters and timers.

5.3.1 PLSQL_PROFILER_RUNS

The PLSQL_PROFILER_RUNS table contains the following columns:

Column	Data Type	Description
runid	INTEGER (NOT NULL)	Unique identifier (plsql_profiler_runnumber)
related_run	INTEGER	The runid of a related run.
run_owner	TEXT	The role that recorded the profiling session.
run_date	TIMESTAMP WITHOUT TIME ZONE	The profiling session start time.
run_comment	TEXT	User comments relevant to this run
run_total_time	BIGINT	Run time (in nanoseconds)
run_system_info	TEXT	Currently Unused
run_comment1	TEXT	Additional user comments
spare1	TEXT	Currently Unused

5.3.2 PLSQL_PROFILER_UNITS

The PLSQL_PROFILER_UNITS table contains the following columns:

Column	Data Type	Description
runid	INTEGER	Unique identifier (plsql_profiler_runnumber)
unit_number	OID	Corresponds to the OID of the row in the pg_proc table that identifies the unit.
unit_type	TEXT	PL/SQL function, procedure, trigger or anonymous block
unit_owner	TEXT	The identity of the role that owns the unit.
unit_name	TEXT	The complete signature of the unit.
unit_timestamp	TIMESTAMP WITHOUT TIME ZONE	Creation date of the unit (currently NULL).
total_time	BIGINT	Time spent within the unit (in nanoseconds)
spare1	BIGINT	Currently Unused
spare2	BIGINT	Currently Unused

5.3.3 PLSQL_PROFILER_DATA

The PLSQL_PROFILER_DATA view contains the following columns:

Column	Data Type	Description
runid	INTEGER	Unique identifier (plsql_profiler_runnumber)
unit_number	OID	Object ID of the unit that contains the current line.
line#	INTEGER	Current line number of the profiled workload.
total_occur	BIGINT	The number of times that the line was executed.
total_time	DOUBLE PRECISION	The amount of time spent executing the line.
min_time	DOUBLE PRECISION	The minimum execution time for the line.
max_time	DOUBLE PRECISION	The maximum execution time for the line.
spare1	NUMBER	Currently Unused
spare2	NUMBER	Currently Unused
spare3	NUMBER	Currently Unused
spare4	NUMBER	Currently Unused

5.3.4 PLSQL_PROFILER_RAWDATA

The PLSQL_PROFILER_RAWDATA table contains the statistical information that is found in the PLSQL_PROFILER_DATA view, as well as the performance statistics returned by the DRITA counters and timers.

Column	Data Type	Description
runid	INTEGER	The run identifier (plsql_profiler_runnumber).
sourcecode	TEXT	The individual line of profiled code.
func_oid	OID	Object ID of the unit that contains the current line.
line_number	INTEGER	Current line number of the profiled workload.
exec_count	BIGINT	The number of times that the line was executed.
time_total	DOUBLE PRECISION	The amount of time spent executing the line.
time_shortest	DOUBLE PRECISION	The minimum execution time for the line.
time_longest	DOUBLE PRECISION	The maximum execution time for the line.
tuples_returned	BIGINT	Currently Unused
num_scans	BIGINT	Currently Unused
tuples_fetched	BIGINT	Currently Unused
tuples_inserted	BIGINT	Currently Unused
tuples_updated	BIGINT	Currently Unused
tuples_deleted	BIGINT	Currently Unused
blocks_fetched	BIGINT	Currently Unused
blocks_hit	BIGINT	Currently Unused
wal_write	BIGINT	The server has waited for a write to the write-ahead log buffer (expect this value to be high).
wal_flush	BIGINT	The server has waited for the write-ahead log to flush to disk.
wal_file_sync	BIGINT	The server has waited for the write-ahead log to sync to disk (related to the wal_sync_method parameter which, by default, is 'fsync' - better performance can be gained by changing this

		parameter to open_sync).
buffer_free_list_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the list of free buffers (in shared memory).
shmem_index_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the shared-memory map.
oid_gen_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the next available OID (object ID).
xid_gen_lock_acquire	bigint	The server has waited for the short-term lock that synchronizes access to the next available transaction ID.
proc_array_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the process array
sinval_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the cache invalidation state.
freespace_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the freespace map.
wal_insert_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes write access to the write-ahead log. A high number may indicate that WAL buffers are sized too small.
wal_write_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes write-ahead log flushes.
control_file_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes write access to the control file (this should usually be a low number).
checkpoint_lock_acquire	BIGINT	A server process has waited for the short-term lock that prevents simultaneous checkpoints.
clog_control_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the commit log.
subtrans_control_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the subtransaction log.
multi_xact_gen_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the next available multi-transaction ID (when a SELECT...FOR SHARE statement executes).
multi_xact_offset_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the multi-transaction offset file (when a SELECT...FOR SHARE statement executes).
multi_xact_member_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the multi-transaction member file (when a SELECT...FOR SHARE statement executes).
rel_cache_init_lock_acquire	BIGINT	The server has waited for the short-term lock that prevents simultaneous relation-cache loads/unloads.
bgwriter_communication_lock_acquire	BIGINT	The bgwriter (background writer) process has waited for the short-term lock that synchronizes messages between the bgwriter and a backend process.
two_phase_state_lock_acquire	BIGINT	The server has waited for the short-term lock that

re		synchronizes access to the list of prepared transactions.
tablespace_create_lock_acquire	BIGINT	The server has waited for the short-term lock that prevents simultaneous CREATE TABLESPACE or DROP TABLESPACE commands.
btree_vacuum_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the next available vacuum cycle ID.
add_in_shmem_lock_acquire	BIGINT	Currently Unused
autovacuum_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the shared autovacuum state.
autovacuum_schedule_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the autovacuum schedule.
syncscan_lock_acquire	BIGINT	The server has waited for the short-term lock that coordinates synchronous scans.
icache_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to InfiniteCache state
breakpoint_lock_acquire	BIGINT	The server has waited for the short-term lock that synchronizes access to the debugger breakpoint list.
lwlock_acquire	BIGINT	The server has waited for a short-term lock that has not been described elsewhere in this section.
db_file_read	BIGINT	A server process has waited for the completion of a read (from disk).
db_file_write	BIGINT	A server process has waited for the completion of a write (to disk).
db_file_sync	BIGINT	A server process has waited for the operating system to flush all changes to disk.
db_file_extend	BIGINT	A server process has waited for the operating system while adding a new page to the end of a file.
sql_parse	BIGINT	Currently Unused
query_plan	BIGINT	The server has generated a query plan.
infinitemcache_read	BIGINT	The server has waited for an Infinite Cache read request.
infinitemcache_write	BIGINT	The server has waited for an Infinite Cache write request.
wal_write_time	BIGINT	The amount of time that the server has waited for a write to the write-ahead log buffer (expect this value to be high).
wal_flush_time	BIGINT	The amount of time that the server has waited for the write-ahead log to flush to disk.
wal_file_sync_time	BIGINT	The amount of time that the server has waited for the write-ahead log to sync to disk (related to the wal_sync_method parameter which, by default, is 'fsync' - better performance can be gained by changing this parameter to open_sync).
buffer_free_list_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the list of free buffers (in shared memory).
shmem_index_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to

		the shared-memory map.
oid_gen_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the next available OID (object ID).
xid_gen_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the next available transaction ID.
proc_array_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the process array.
sinval_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the cache invalidation state.
freespace_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the freespace map.
wal_insert_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes write access to the write-ahead log. A high number may indicate that WAL buffers are sized too small.
wal_write_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes write-ahead log flushes.
control_file_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes write access to the control file (this should usually be a low number).
checkpoint_lock_acquire_time	BIGINT	The amount of time that the server process has waited for the short-term lock that prevents simultaneous checkpoints.
clog_control_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the commit log.
subtrans_control_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the subtransaction log.
multi_xact_gen_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the next available multi-transaction ID (when a SELECT...FOR SHARE statement executes).
multi_xact_offset_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the multi-transaction offset file (when a SELECT...FOR SHARE statement executes).
multi_xact_member_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the multi-transaction member file (when a SELECT...FOR SHARE statement executes).
rel_cache_init_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that prevents simultaneous relation-cache loads/unloads.
bgwriter_communication_lock_acquire_time	BIGINT	The amount of time that the bgwriter (background writer) process has waited for the short-term lock that synchronizes messages

		between the bgwriter and a backend process.
two_phase_state_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the list of prepared transactions.
tablespace_create_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that prevents simultaneous CREATE TABLESPACE or DROP TABLESPACE commands.
btree_vacuum_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the next available vacuum cycle ID.
add_in_shmem_lock_acquire_time	BIGINT	Obsolete/unused
autovacuum_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the shared autovacuum state.
autovacuum_schedule_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the autovacuum schedule.
syncscan_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that coordinates synchronous scans.
icache_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to InfiniteCache state
breakpoint_lock_acquire_time	BIGINT	The amount of time that the server has waited for the short-term lock that synchronizes access to the debugger breakpoint list.
lwlock_acquire_time	BIGINT	The amount of time that the server has waited for a short-term lock that has not been described elsewhere in this section.
db_file_read_time	BIGINT	The amount of time that the server process has waited for the completion of a read (from disk).
db_file_write_time	BIGINT	The amount of time that the server process has waited for the completion of a write (to disk).
db_file_sync_time	BIGINT	The amount of time that the server process has waited for the operating system to flush all changes to disk.
db_file_extend_time	BIGINT	The amount of time that the server process has waited for the operating system while adding a new page to the end of a file.
sql_parse_time	BIGINT	The amount of time that the server has parsed a SQL statement.
query_plan_time	BIGINT	The amount of time that the server has computed the execution plan for a SQL statement.
infinitecache_read_time	BIGINT	The amount of time that the server has waited for an Infinite Cache read request.
infinitecache_write_time	BIGINT	The amount of time that the server has waited for an Infinite Cache write request.
totalwaits	BIGINT	The total number of event waits.
totalwaittime	BIGINT	The total time spent waiting for an event.

6 Index Advisor

The Index Advisor utility helps determine which columns you should index to improve performance in a given workload. Index Advisor considers B-tree (single-column or composite) index types, and does not identify other index types (GIN, GiST, Hash) that may improve performance. Index Advisor is installed with Postgres Plus Advanced Server.

Index Advisor works with Advanced Server's query planner by creating *hypothetical indexes* that the query planner uses to calculate execution costs as if such indexes were available. Index Advisor identifies the indexes by analyzing SQL queries supplied in the workload.

There are three ways to use Index Advisor to analyze SQL queries:

- Invoke the Index Advisor utility program, supplying a text file containing the SQL queries that you wish to analyze; Index Advisor will generate a text file with `CREATE INDEX` statements for the recommended indexes.
- Provide queries at the EDB-PSQL command line that you want Index Advisor to analyze.
- Access Index Advisor through the Postgres Enterprise Manager client. When accessed via the PEM client, Index Advisor works with SQL Profiler, providing indexing recommendations on code captured in SQL traces. For more information about using SQL Profiler and Index Advisor with PEM, please see Section 8.4 of the *PEM Getting Started Guide*, available from the EnterpriseDB website at:

<http://www.enterprisedb.com/products-services-training/products/postgres-enterprise-manager>

Index Advisor will attempt to make indexing recommendations on `INSERT`, `UPDATE`, `DELETE` and `SELECT` statements. When invoking Index Advisor, you supply the workload in the form of a set of queries (if you are providing the command in an SQL file) or an `EXPLAIN` statement (if you are specifying the SQL statement at the psql command line). Index Advisor displays the query plan and estimated execution cost for the supplied query, but does not actually execute the query.

During the analysis, Index Advisor compares the query execution costs with and without hypothetical indexes. If the execution cost using a hypothetical index is less than the execution cost without it, both plans are reported in the `EXPLAIN` statement output, metrics that quantify the improvement are calculated, and Index Advisor generates the `CREATE INDEX` statement needed to create the index.

If no hypothetical index can be found that reduces the execution cost, Index Advisor displays only the original query plan output of the `EXPLAIN` statement.

Index Advisor does not actually create indexes on the tables. Use the `CREATE INDEX` statements supplied by Index Advisor to add any recommended indexes to your tables.

A script supplied with Advanced Server creates the table in which Index Advisor stores the indexing recommendations generated by the analysis; the script also creates a function and a view of the table to simplify the retrieval and interpretation of the results.

If you choose to forego running the script, Index Advisor will log recommendations in a temporary table that is available only for the duration of the Index Advisor session.

6.1 Index Advisor Components

The Index Advisor shared library interacts with the query planner to make indexing recommendations. The Postgres Plus Advanced Server installer creates the following shared library in the `libdir` subdirectory of your Postgres Plus Advanced Server home directory:

On Linux:

```
index_advisor.so
```

On Windows:

```
index_advisor.dll
```

Please note that libraries in the `libdir` directory can only be loaded by a superuser. A database administrator can allow a non-superuser to use Index Advisor by manually copying the Index Advisor file from the `libdir` directory into the `libdir/plugins` directory (under your Advanced Server home directory). Only a trusted non-superuser should be allowed access to the plugin; this is an unsafe practice in a production environment.

The installer also creates the Index Advisor utility program and setup script:

```
pg_advise_index
```

`pg_advise_index` is a utility program that reads a user-supplied input file containing SQL queries and produces a text file containing `CREATE INDEX` statements that can be used to create the indexes recommended by the Index Advisor. The `pg_advise_index` program is located in the `bin` subdirectory of the Postgres Plus Advanced Server home directory.

`index_advisor.sql`

`index_advisor.sql` is a script that creates a permanent IndexAdvisor log table along with a function and view to facilitate reporting of recommendations from the log table. The script is located in the `share/contrib` subdirectory of the Postgres Plus Advanced Server directory.

The `index_advisor.sql` script creates the `index_advisor_log` table, the `show_index_recommendations()` function and the `index_recommendations` view. These database objects must be created in a schema that is accessible by, and included in the search path of the role that will invoke IndexAdvisor.

`index_advisor_log`

Index Advisor logs indexing recommendations in the `index_advisor_log` table. If Index Advisor does not find the `index_advisor_log` table in the user's search path, IndexAdvisor will store any indexing recommendations in a temporary table of the same name. The temporary table exists only for the duration of the current session.

`show_index_recommendations()`

`show_index_recommendations()` is a PL/pgSQL function that interprets and displays the recommendations made during a specific IndexAdvisor session (as identified by its backend process ID).

`index_recommendations`

Index Advisor creates the `index_recommendations` view based on information stored in the `index_advisor_log` table during a query analysis. The view produces output in the same format as the `show_index_recommendations()` function, but contains IndexAdvisor recommendations for all stored sessions, while the result set returned by the `show_index_recommendations()` function are limited to a specified session.

6.2 Index Advisor Configuration

Index Advisor does not require any configuration to generate recommendations that are available only for the duration of the current session; to store the results of multiple sessions, you must create the `index_advisor_log` table (where Advanced Server will store IndexAdvisor recommendations). To create the `index_advisor_log` table, you must run the `index_advisor.sql` script.

When selecting a storage schema for the IndexAdvisor table, function and view, keep in mind that all users that invoke IndexAdvisor (and query the result set) must have `USAGE` privileges on the schema. The schema must be in the search path of all users that are interacting with the IndexAdvisor.

1. Place the selected schema at the start of your `search_path` parameter. For example, if your search path is currently:

```
search_path=public, accounting
and you want the IndexAdvisor objects to be created in a schema named
advisor, use the command:
SET search_path = advisor, public, accounting;
```

2. Run the `index_advisor.sql` script to create the database objects. If you are running the `psql` client, you can use the command:

```
\i full_pathname/index_advisor.sql
Specify the pathname to the index_advisor.sql script in place of
full_pathname.
```

3. Grant privileges on the `index_advisor_log` table to all IndexAdvisor users; this step is not necessary if the IndexAdvisor user is a superuser, or the owner of these database objects.
 - Grant `SELECT` and `INSERT` privileges on the `index_advisor_log` table to allow a user to invoke IndexAdvisor.
 - Grant `DELETE` privileges on the `index_advisor_log` table to allow the specified user to delete the table contents.
 - Grant `SELECT` privilege on the `index_recommendations` view.

The following example demonstrates the creation of the IndexAdvisor database objects in a schema named `ia`, which will then be accessible to an IndexAdvisor user with user name `ia_user`:

```
$ edb-psql -d edb -U enterprisedb
edb-psql (9.2.0.0)
Type "help" for help.

edb=# CREATE SCHEMA ia;
CREATE SCHEMA
edb=# SET search_path TO ia;
SET
edb=# \i /opt/PostgresPlus/9.2AS/share/contrib/index_advisor.sql
CREATE TABLE
CREATE INDEX
CREATE INDEX
CREATE FUNCTION
```

```
CREATE FUNCTION
CREATE VIEW
edb=# GRANT USAGE ON SCHEMA ia TO ia_user;
GRANT
edb=# GRANT SELECT, INSERT, DELETE ON index_advisor_log TO ia_user;
GRANT
edb=# GRANT SELECT ON index_recommendations TO ia_user;
GRANT
```

While using IndexAdvisor, the specified schema (*ia*) must be included in *ia_user*'s *search_path* parameter.

6.3 Using Index Advisor

When you invoke IndexAdvisor, you must supply a workload; the workload is either a query (specified at the command line), or a file that contains a set of queries (executed by the `pg_advise_index()` function). After analyzing the workload, IndexAdvisor will either store the result set in a temporary table, or in a permanent table. You can review the indexing recommendations generated by IndexAdvisor and use the `CREATE INDEX` statements generated by IndexAdvisor to create the recommended indexes.

Note: You should not run IndexAdvisor in read-only transactions.

The following examples assume that `superuser` `enterprisedb` is the IndexAdvisor user, and the IndexAdvisor database objects have been created in a schema in the *search_path* of `superuser` `enterprisedb`.

The examples in the following sections use the table created with the statement shown below:

```
CREATE TABLE t( a INT, b INT );
INSERT INTO t SELECT s, 99999 - s FROM generate_series(0,99999) AS s;
ANALYZE t;
```

The resulting table contains the following rows:

a	b
0	99999
1	99998
2	99997
3	99996
.	.
99997	2
99998	1
99999	0

6.3.1 Using the pg_advise_index Utility

When invoking the `pg_advise_index` utility, you must include the name of a file that contains the queries that will be executed by `pg_advise_index`; the queries may be on the same line, or on separate lines, but each query must be terminated by a semicolon. Queries within the file should not begin with the `EXPLAIN` keyword.

The following example shows the contents of a sample `workload.sql` file:

```
SELECT * FROM t WHERE a = 500;
SELECT * FROM t WHERE b < 1000;
```

Run the `pg_advise_index` program as shown in the code sample below:

```
$ pg_advise_index -d edb -h localhost -U enterprisedb -s 100M -o advisory.sql
workload.sql
poolsize = 102400 KB
load workload from file 'workload.sql'
Analyzing queries .. done.
size = 2184 KB, benefit = 1684.720000
size = 2184 KB, benefit = 1655.520000
/* 1. t(a): size=2184 KB, benefit=1684.72 */
/* 2. t(b): size=2184 KB, benefit=1655.52 */
/* Total size = 4368KB */
```

In the code sample, the `-d`, `-h`, and `-U` options are `psql` connection options.

`-s`

`-s` is an optional parameter that limits the maximum size of the indexes recommended by Index Advisor. If Index Advisor does not return a result set, `-s` may be set too low.

`-o`

The recommended indexes are written to the file specified after the `-o` option.

The information displayed by the `pg_advise_index` program is logged in the `index_advisor_log` table. In response to the command shown in the example, Index Advisor writes the following `CREATE INDEX` statements to the `advisory.sql` output file

```
create index idx_t_1 on t (a);
create index idx_t_2 on t (b);
```

You can create the recommended indexes at the `psql` command line with the `CREATE INDEX` statements in the file, or create the indexes by executing the `advisory.sql` script.

```
$ edb-psql -d edb -h localhost -U enterprisedb -e -f advisory.sql
create index idx_t_1 on t (a);
CREATE INDEX
create index idx_t_2 on t (b);
CREATE INDEX
```

6.3.2 Using Index Advisor at the psql Command Line

You can use IndexAdvisor to analyze SQL statements entered at the edb-psql (or psql) command line; the following steps detail loading the IndexAdvisor plugin and using Index Advisor:

1. Connect to the server with the edb-psql command line utility, and load the Index Advisor plugin:

```
$ edb-psql -d edb -U enterprisedb
...
edb=# LOAD 'index_advisor';
LOAD
```

2. Use the edb-psql command line to invoke each SQL command that you would like Index Advisor to analyze. IndexAdvisor stores any recommendations for the queries in the index_advisor_log table. If the index_advisor_log table does not exist in the user's search_path, a temporary table is created with the same name. This temporary table exists only for the duration of the user's session.

After loading the IndexAdvisor plugin, IndexAdvisor will analyze all SQL statements and log any indexing recommendations for the duration of the session.

If you would like IndexAdvisor to analyze a query (and make indexing recommendations) without actually executing the query, preface the SQL statement with the EXPLAIN keyword.

If you do not preface the statement with the EXPLAIN keyword, IndexAdvisor will analyze the statement while the statement executes, writing the indexing recommendations to the index_advisor_log table for later review.

In the example that follows, the EXPLAIN statement displays the normal query plan, followed by the query plan of the same query, if the query were using the recommended hypothetical index:

```
edb=# EXPLAIN SELECT * FROM t WHERE a < 10000;
               QUERY PLAN
-----
Seq Scan on t  (cost=0.00..1693.00 rows=10105 width=8)
  Filter: (a < 10000)
Result  (cost=0.00..337.10 rows=10105 width=8)
  One-Time Filter: '==[ HYPOTHETICAL PLAN ]==':text
```

```

-> Index Scan using "<hypothetical-index>:1" on t
    (cost=0.00..337.10 rows=10105 width=8)
    Index Cond: (a < 10000)
(6 rows)

edb=# EXPLAIN SELECT * FROM t WHERE a = 100;
                                QUERY PLAN
-----
Seq Scan on t  (cost=0.00..1693.00 rows=1 width=8)
  Filter: (a = 100)
Result  (cost=0.00..8.28 rows=1 width=8)
  One-Time Filter: '==[ HYPOTHETICAL PLAN ]==':text
  -> Index Scan using "<hypothetical-index>:3" on t
      (cost=0.00..8.28 rows=1 width=8)
      Index Cond: (a = 100)
(6 rows)

```

For information about reviewing the recommended queries, see Section 8.4, *Reviewing the Index Advisor Recommendations*.

After loading the IndexAdvisor plugin, the default value of `index_advisor.enabled` is on. The IndexAdvisor plugin must be loaded to use a `SET` or `SHOW` command to display the current value of `index_advisor.enabled`.

You can use the `index_advisor.enabled` parameter to temporarily disable Index Advisor without interrupting the psql session:

```

edb=# SET index_advisor.enabled TO off;
SET

```

To enable IndexAdvisor, set the parameter to on:

```

edb=# SET index_advisor.enabled TO on;
SET

```

6.4 Reviewing the Index Advisor Recommendations

There are several ways to review the index recommendations generated by Index Advisor. You can:

- Query the `index_advisor_log` table.
- Run the `show_index_recommendations` function.
- Query the `index_recommendations` view.

6.4.1 Using the show_index_recommendations() Function

To review the recommendations of the Index Advisor utility using the `show_index_recommendations()` function, call the function, specifying the process ID of the session:

```
SELECT show_index_recommendations( pid );
```

Where `pid` is the process ID of the current session. If you do not know the process ID of your current session, passing a value of `NULL` will also return a result set for the current session.

The following code fragment shows an example of a row in a result set:

```
edb=# SELECT show_index_recommendations(null);
          show_index_recommendations
-----
create index idx_t_a on t(a);/* size: 2184 KB, benefit: 3040.62,
gain: 1.39222666981456 */
(1 row)
```

In the example, `create index idx_t_a on t(a)` is the SQL statement needed to create the index suggested by Index Advisor. Each row in the result set shows:

- The command required to create the recommended index.
- The maximum estimated size of the index.
- The calculated benefit of using the index.
- The estimated gain that will result from implementing the index.

You can display the results of all Index Advisor sessions from the following view:

```
SELECT * FROM index_recommendations;
```

6.4.2 Querying the index_advisor_log Table

Index Advisor stores indexing recommendations in a table named `index_advisor_log`. Each row in the `index_advisor_log` table contains the result of a query where Index Advisor determines it can recommend a hypothetical index to reduce the execution cost of that query.

Column	Type	Description
relid	oid	OID of the base table for the index
relname	name	Name of the base table for the index
attrs	integer[]	Recommended index columns (identified by column number)
benefit	real	Calculated benefit of the index for this query

index_size	integer	Estimated index size in disk-pages
backend_pid	integer	Process ID of the process generating this recommendation
timestamp	timestamp	Date/Time when the recommendation was generated

You can query the `index_advisor_log` table at the `psql` command line. The following example shows the `index_advisor_log` table entries resulting from two Index Advisor sessions. Each session contains two queries, and can be identified (in the table below) by a different `backend_pid` value. For each session, Index Advisor generated two index recommendations.

```
edb=# SELECT * FROM index_advisor_log;
   reloid | relname | attrs | benefit | index_size | backend_pid |
timestamp
-----+-----+-----+-----+-----+-----+-----
16651 | t       | {1}   | 1684.72 | 2184 | 3442 | 22-MAR-11
16:44:32.712638 -04:00
16651 | t       | {2}   | 1655.52 | 2184 | 3442 | 22-MAR-11
16:44:32.759436 -04:00
16651 | t       | {1}   | 1355.9  | 2184 | 3506 | 22-MAR-11
16:48:29.317016 -04:00
16651 | t       | {1}   | 1684.72 | 2184 | 3506 | 22-MAR-11
16:51:45.927906 -04:00
(4 rows)
```

Index Advisor added the first two rows to the table after analyzing the following two queries executed by the `pg_advise_index` utility:

```
SELECT * FROM t WHERE a = 500;
SELECT * FROM t WHERE b < 1000;
```

The value of 3442 in column `backend_pid` identifies these results as coming from the session with process ID 3442.

The value of 1 in column `attrs` in the first row indicates that the hypothetical index is on the first column of the table (column `a` of table `t`).

The value of 2 in column `attrs` in the second row indicates that the hypothetical index is on the second column of the table (column `b` of table `t`).

Index Advisor added the last two rows to the table after analyzing the following two queries (executed at the `psql` command line):

```
edb=# EXPLAIN SELECT * FROM t WHERE a < 10000;
               QUERY PLAN
-----
Seq Scan on t (cost=0.00..1693.00 rows=10105 width=8)
  Filter: (a < 10000)
Result (cost=0.00..337.10 rows=10105 width=8)
  One-Time Filter: '===[ HYPOTHETICAL PLAN ]==':text
```

```

-> Index Scan using "<hypothetical-index>:1" on t (cost=0.00..337.10
rows=10105 width=8)
      Index Cond: (a < 10000)
(6 rows)

edb=# EXPLAIN SELECT * FROM t WHERE a = 100;
               QUERY PLAN
-----
Seq Scan on t (cost=0.00..1693.00 rows=1 width=8)
  Filter: (a = 100)
Result (cost=0.00..8.28 rows=1 width=8)
  One-Time Filter: '==[ HYPOTHETICAL PLAN ]==':text
-> Index Scan using "<hypothetical-index>:3" on t (cost=0.00..8.28
rows=1 width=8)
      Index Cond: (a = 100)
(6 rows)

```

The values in the `benefit` column of the `index_advisor_log` table are calculated using the following formula:

$$\text{benefit} = (\text{normal execution cost}) - (\text{execution cost with hypothetical index})$$

The value of the `benefit` column for the last row of the `index_advisor_log` table (shown in the example) is calculated using the query plan for the following SQL statement:

```
EXPLAIN SELECT * FROM t WHERE a = 100;
```

The execution costs of the different execution plans are evaluated and compared:

$$\text{benefit} = (\text{Seq Scan on t cost}) - (\text{Index Scan using <hypothetical-index>})$$

and the benefit is added to the table:

$$\begin{aligned} \text{benefit} &= 1693.00 - 8.28 \\ \text{benefit} &= 1684.72 \end{aligned}$$

You can delete rows from the `index_advisor_log` table when you no longer have the need to review the results of the queries stored in the row.

6.4.3 Querying the `index_recommendations` View

The `index_recommendations` view contains the calculated metrics and the `CREATE INDEX` statements to create the recommended indexes for all sessions whose results are currently in the `index_advisor_log` table. You can display the results of all stored Index Advisor sessions by querying the `index_recommendations` view as shown below:

```
SELECT * FROM index_recommendations;
```

Using the example shown in the previous section (*Querying the index_advisor_log Table*), the `index_recommendations` view displays the following:

```
edb=# SELECT * FROM index_recommendations;
 backend_pid | show_index_recommendations
-----+-----
3442 | create index idx_t_a on t(a);/* size: 2184 KB, benefit:
1684.72, gain: 0.771392654586624 */
3442 | create index idx_t_b on t(b);/* size: 2184 KB, benefit:
1655.52, gain: 0.758021539820856 */
3506 | create index idx_t_a on t(a);/* size: 2184 KB, benefit:
3040.62, gain: 1.39222666981456 */
(3 rows)
```

Within each session, the results of all queries that benefit from the same recommended index are combined to produce one set of metrics per recommended index, reflected in the fields named `benefit` and `gain`.

The formulas for the fields are as follows:

```
size      = MAX(index size of all queries)
benefit   = SUM(benefit of each query)
gain      = SUM(benefit of each query) / MAX(index size of all queries)
```

So for example, using the following query results from the process with a `backend_pid` of 3506:

```
reloid | relname | attrs | benefit | index_size | backend_pid |
timestamp
-----+-----+-----+-----+-----+-----+-----
16651 | t       | {1}   | 1355.9  | 2184       | 3506 | 22-MAR-11
16:48:29.317016 -04:00
16651 | t       | {1}   | 1684.72 | 2184       | 3506 | 22-MAR-11
16:51:45.927906 -04:00
```

The metrics displayed from the `index_recommendations` view for `backend_pid` 3506 are:

```
 backend_pid | show_index_recommendations
-----+-----
3506 | create index idx_t_a on t(a);/* size: 2184 KB, benefit:
3040.62, gain: 1.39222666981456 */
```

The metrics from the view are calculated as follows:

```
benefit = (benefit from 1st query) + (benefit from 2nd
query)
benefit = 1355.9 + 1684.72
```

```
benefit = 3040.62
```

and

```
gain = ((benefit from 1st query) + (benefit from 2nd
query)) / MAX(index size of all queries)
gain = (1355.9 + 1684.72) / MAX(2184, 2184)
gain = 3040.62 / 2184
gain = 1.39223
```

The gain metric is useful when comparing the relative advantage of the different recommended indexes derived during a given session. The larger the gain value, the better the cost effectiveness derived from the index weighed against the possible disk space consumption of the index.

6.5 Limitations

Index Advisor does not consider IndexOnly scans; it does consider Indexscans when making recommendations.

Index Advisor ignores any computations found in the `WHERE` clause. Effectively, the index field in the recommendations will not be any kind of expression; the field will be a simple column name.

Index Advisor does not consider inheritance when recommending hypothetical indexes. If a query references a parent table, Index Advisor does not make any index recommendations on child tables.

Restoration of a `pg_dump` backup file that includes the `index_advisor_log` table or any tables for which indexing recommendations were made and stored in the `index_advisor_log` table, may result in "broken links" between the `index_advisor_log` table and the restored tables referenced by rows in the `index_advisor_log` table because of changes in object identifiers (OIDs).

If it is necessary to display the recommendations made prior to the backup, you can replace the old OIDs in the `reloid` column of the `index_advisor_log` table with the new OIDs of the referenced tables using the `SQL UPDATE` statement:

```
UPDATE index_advisor_log SET reloid = new_oid WHERE reloid
= old_oid;
```

7 Other Performance Features

This chapter provides a brief summary of other performance related features of Postgres Plus Advanced Server.

7.1 SQL Profiler

Inefficient SQL code is one of, if not the leading cause of database performance problems. The challenge for database administrators and developers is locating and then optimizing this code in large, complex systems.

SQL Profiler helps you locate and optimize poorly running SQL code.

Specific features and benefits of SQL Profiler include the following:

- **On-Demand Traces.** You can capture SQL traces at any time by manually setting up your parameters and starting the trace.
- **Scheduled Traces.** For inconvenient times, you can also specify your trace parameters and schedule them to run at some later time.
- **Save Traces.** Execute your traces and save them for later review.
- **Trace Filters.** Selectively filter SQL captures by database and by user, or capture every SQL statement sent by all users against all databases.
- **Trace Output Analyzer.** A graphical table lets you quickly sort and filter queries by duration or statement, and a graphical or text based `EXPLAIN` plan lays out your query paths and joins.
- **Index Advisor Integration.** Once you have found your slow queries and optimized them, you can also let the Index Advisor recommend the creation of underlying table indices to further improve performance.

For more information about SQL Profiler and Postgres Enterprise Manager, visit the EnterpriseDB website at:

<http://www.enterprisedb.com/postgres-enterprise-manager>

7.2 Query Optimization Hints

The Advanced Server query planner performs the task of determining how the result set should be produced for `DELETE`, `SELECT`, and `UPDATE` SQL commands. The query planner uses cost based optimization to determine the least cost plan.

There may be cases where you want to force the query planner to either utilize (or not to utilize) a certain access method to determine the plan (for example, use a sequential scan instead of an index scan).

The *query hints* feature (also called *optimizer hints*) of Advanced Server allows you to embed these directives within the SQL command to force the query planner to use a certain access method.

Query hints are typically used with the `EXPLAIN` command so you can see the estimated costs associated with the generated plan based on the query hint you supplied.

This technique allows you to perform a more detailed cost comparison of how certain queries are expected to perform.

For more information about optimizer hints, see the Oracle Compatibility Developer's Guide, available from the EnterpriseDB website at:

<http://www.enterprisedb.com/documentation>

7.3 Hi-Speed Bulk Loader

Hi-speed bulk loading is provided by the Postgres Plus Advanced Server *EDB*Loader* feature. EDB*Loader is a command line utility that loads data from an input source, typically a file, into one or more tables using directives compatible with Oracle SQL*Loader.

EDB*Loader provides the following features and benefits:

- Three data loading methods providing various levels of performance – conventional path load, direct path load, and parallel direct path load
- *Conventional path load* for full insert processing, which includes all integrity checks
- *Direct path load* for faster performance by writing data directly to database pages bypassing insert processing overhead of a conventional path load
- *Parallel direct path load* for even faster performance by distributing the loading process over multiple, parallel sessions
- Oracle SQL*Loader compatible syntax for control file directives
- Input data with delimiter-separated or fixed-width fields
- Bad file for collecting rejected records
- Loading of multiple target tables
- Discard file for collecting records that do not meet the selection criteria of any target table
- Log file for recording the EDB*Loader session and any error messages
- Data loading from standard input and remote loading, particularly useful for large data sources on remote hosts

For more information about EDB*Loader, see the Oracle Compatibility Developer's Guide, available from the EnterpriseDB website at:

<http://www.enterprisedb.com/documentation>

7.4 Bulk Collect Fetch and Binding

SQL statements that return a result set consisting of a large number of rows may not be operating as efficiently as possible due to the constant context switching that must occur between the database server and the client in order to transfer the entire result set.

Performance can be improved in this scenario by using a collection (that is, an array) to gather the entire result set in memory, which the client can then access. This is accomplished by creating a collection and then using the SQL statement with the `BULK COLLECT` clause to gather the result set in the collection.

This performance enhancement can be applied to the following types of SQL statements:

- `SELECT INTO`
- `FETCH INTO`
- `EXECUTE IMMEDIATE INTO`
- `DELETE RETURNING INTO`
- `INSERT RETURNING INTO`
- `UPDATE RETURNING INTO`

For more information, see the Oracle Compatibility Developer's Guide, available from the EnterpriseDB website at:

<http://www.enterprisedb.com/documentation>

7.5 Multi-Threaded Replication

The *Postgres Plus xDB Replication Server* is a comprehensive, yet easy-to-use system for replicating tables between database servers. Not only can data be replicated using Postgres Plus Advanced Server databases, but tables from Oracle, Microsoft® SQL Server®, or PostgreSQL® databases can be replicated to Advanced Server and vice versa.

xDB Replication Server provides the following benefits and features:

- Replicate Oracle, SQL Server, or PostgreSQL tables and views to Advanced Server and vice versa

- Distributed multi publication/subscription architecture
- Synchronize data across geographies
- Multi-threaded replication support
- Snapshot and continuous replication modes
- Row filtering
- Flexible replication scheduling
- Cascading replication support
- Replication history viewer
- Management via a graphical replication console or a command line driven interface

xDB Replication Server takes advantage of the system architecture of the host database server. Replication can be done in parallel mode with multiple threads on a multi-CPU or core architecture.

For more information, see the Postgres Plus xDB Replication Server User's Guide, available from the EnterpriseDB website at:

<http://www.enterprisedb.com/docs/en/9.1/repguide/Contents.htm>